



# Introducing Elastic Search, Logstash and Kibana

---

David THIBAU – 2022

[david.thibau@gmail.com](mailto:david.thibau@gmail.com)



# Agenda

---

- Introduction
  - ELK proposition and its use cases
  - Components of the stack
  - Distributions : Community versus X-Pack
  - Built-in features
  - Deployment alternatives
  - Core concepts of clustering
- Data ingestion
  - Beats concepts, provided beats by Elastic,
  - Logstash Concepts: pipelines, input, filters and output plugins
  - Elastic search pipelines
- Search engine capabilities
  - Inverted index and distributed search
  - Supported Data Types
  - Analyzers, predined and custom
  - DSL syntax for queries
  - Full text and relevance score
  - Advanced search
  - Aggregations or faceting
  - Geo queries
- Data Visualization with Kibana
  - What is Nearly-Real-time analysis?
  - Discover, Visualize and Dashboards
  - Basic visualisation
  - Time series
  - Location analysis
  - Data relationships with graph
  - Machine Learning



# Introduction

---

## **ELK proposition and use cases**

Distributions : Community vs X-Pack

Built-in features

Deployment alternatives

Core concepts of clustering



# Introduction

---

The elastic company provides the ***ElasticStack suite*** in 2 editions (OpenSource and Enterprise)

The suite addresses 2 aspects, which can be applied in a BigData environment :

- **Near Real-time search**

Permanently indexed data is available for search, the latency due to indexing is negligible (Near Real Time)

- **Near Real-time data analysis.**

Data is aggregated in real time to provide visualizations and dashboards to extract insights

The core component of this stack is ***ElasticSearch***



# *ElasticSearch*

---

ElasticSearch is a server offering a REST API<sup>1</sup> :

- To store and index data  
(Office documents, tweets, log files, performance metrics, ...)
- To perform search queries (structured, full-text, natural language, geolocation)
- Aggregate data in order to calculate metrics

*1. There is no User Interface embeded in the product Elastic Search*



# Features

---

- ✓ Massively **distributed and scalable** architecture in volume and load  
=> Big Data, remarkable performance
- ✓ **High availability** through replication
- ✓ **Muti-tenancy** or multi-index. The document base contains several indexes whose life cycles are completely independent
- ✓ Based on the **Lucene** reference library  
=> Full-text search, consideration of languages, natural language, etc.
- ✓ **Flexibility** : Structured storage of documents but without prior schema, possibility of adding fields to an existing schema
- ✓ Very complete **RESTFul API**
- ✓ **Open Source**



# Apache Lucene

---

- Apache project Started in 2000, used in many products
- The "low layer" of ELS: a Java library for writing and searching **index files**:
  - An index contains **documents**
  - A document contains **fields**
  - A field consists of **terms**

ELS brings scalability, a REST API, simplicity and aggregation features





# ELS vs SolR

The Apache Foundation offers **SolR** which is very close to ELS in terms of functionality

	SOLR	ELS
<b>Installation &amp; Configuration</b>	Not simple, Very detailed documentation	Simple and intuitive
<b>Indexation/ Search</b>	Natural language oriented	Text and other data types for aggregations
<b>Scalability</b>	Clustering via ZooKeeper and SolRCloud	Natively cluster
<b>Community</b>	Important but stagnant	Has exploded
<b>Documentation</b>	Very complete but very technical	Very complete, many examples





# Elastic Stack

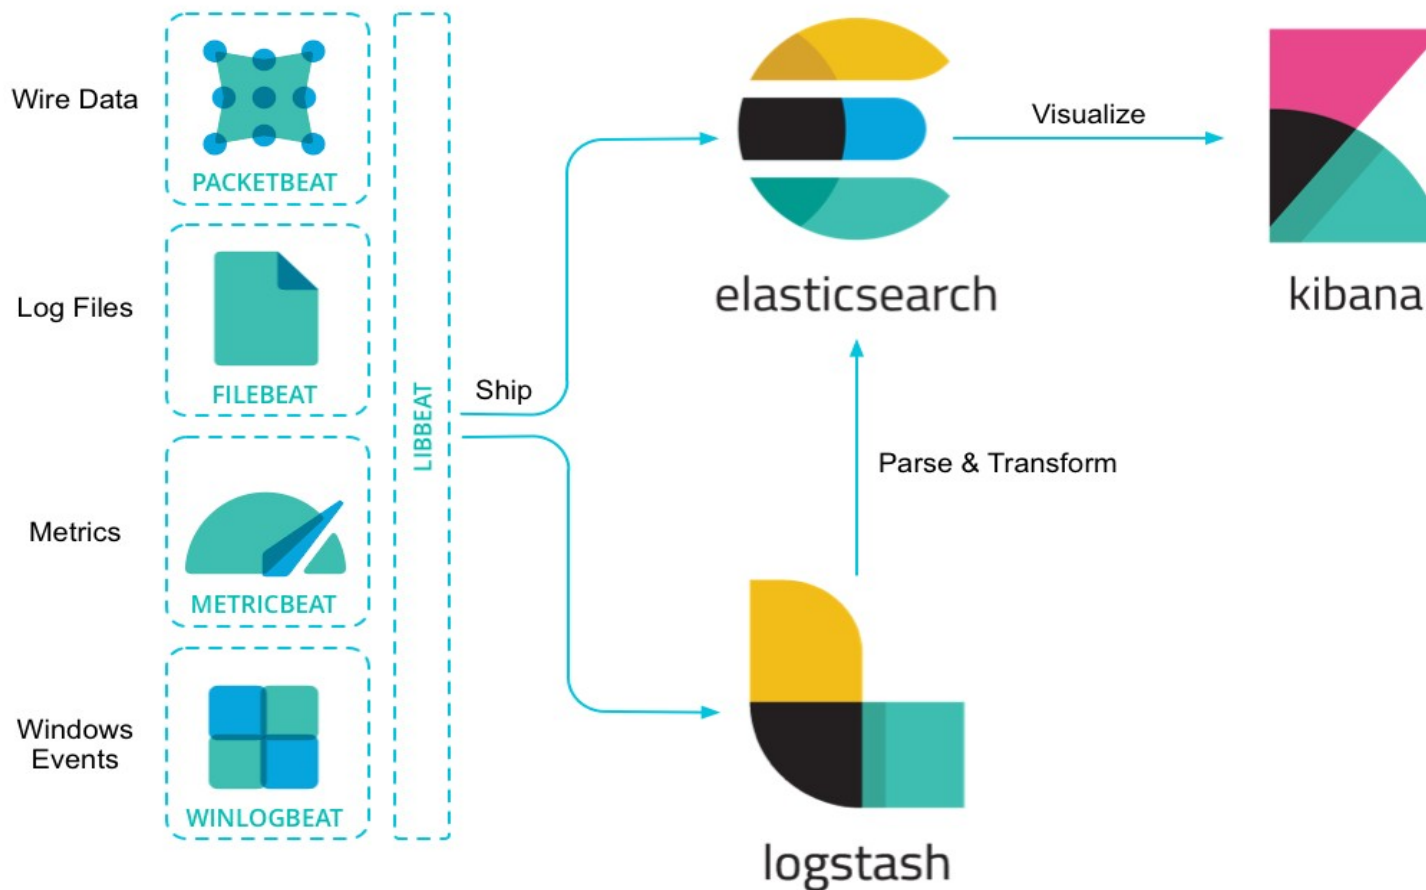
---

The suite Elastic Stack is dedicated to Real Time Analysis in the context of BigData.

It consist of :

- ***Logstash / Beat*** : Ingestion of data.  
Extract and transform
- ***Elastic Search*** : Storage, Indexing and Search
- ***Kibana*** : Data visualization

# Architecture





# Beats

---

Beats are little programs installed on target devices of server

They periodically sends data :

- Directly to ElasticSearch
- Or to a logstash pipelines



# Logstash

---

Logstash is a tool which takes as input events which can come from different sources

It transform, enrich filter these events

Then, it writes the transformed event in different destinations but mainly in ElasticSearc indexes

# Clients Alternatives for ElasticSearch



---

- **Classical REST clients:**

*curl, Postman, SOAPUI, ...*

- **Libraries** provided by Elastic :

*Java, Javascript, Python, Groovy, Php, .NET, Perl*

Developers can easily develop programs which integrates ElasticSearch

- **Kibana** is a generic client for managing cluster, provide dashboards, develop queries and many more



# Kibana

---

Kibana is an application for :

- Data managers
- Operators of the stack
- Developers of application which integrates ElasticSearch

It has lot of menus and also lot of marketing contents



# Introduction

---

ELK proposition and use cases  
**Distributions : Community vs X-Pack**

Built-in features  
Deployment alternatives  
Core concepts of clustering



# Releases

---

Releases are available for download in different formats :

- Archive ZIP, TAR,
- Packages DEB or RPM
- Docker Images

## Versions

- 8.0.0 : February 2022
- 7.x : Avril 2019 - Continuing
- 6.x : November 2017 to December 2019
- 5.0.0 : October 2016
- 2.4.1 : September 2016
- 2.4.0 : August 2016





# Related Offers

---

The company Elastic also offers:

- ***X-Pack*** : commercial and enterprise features
  - Security, Monitoring (before release 8.x)
  - Alerts, Machine Learning
  - Application Performance Management
  - Reporting and scheduling
- ***Elastic Cloud*** : Delegate operations of the cluster to the elastic company or deploy the solution in a private cloud



# Introduction

---

ELK proposition and use cases  
Distributions : Community vs X-Pack  
**Built-in features**  
Deployment alternatives  
Core concepts of clustering



# Solutions out-of-the-box

---

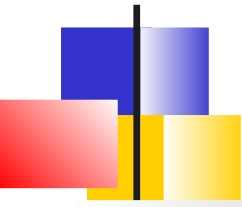
***Enterprise Search*** : Define sources of content (web, other..), index content and provides search capabilities

***Elastic Observability*** : Monitor CPU and memory consumption, aggregate logs, inspect network traffic, Application Performance Management .

Machine Learning to detect anomalies and anticipate future

***Elastic Security*** : Detect and respond to threats

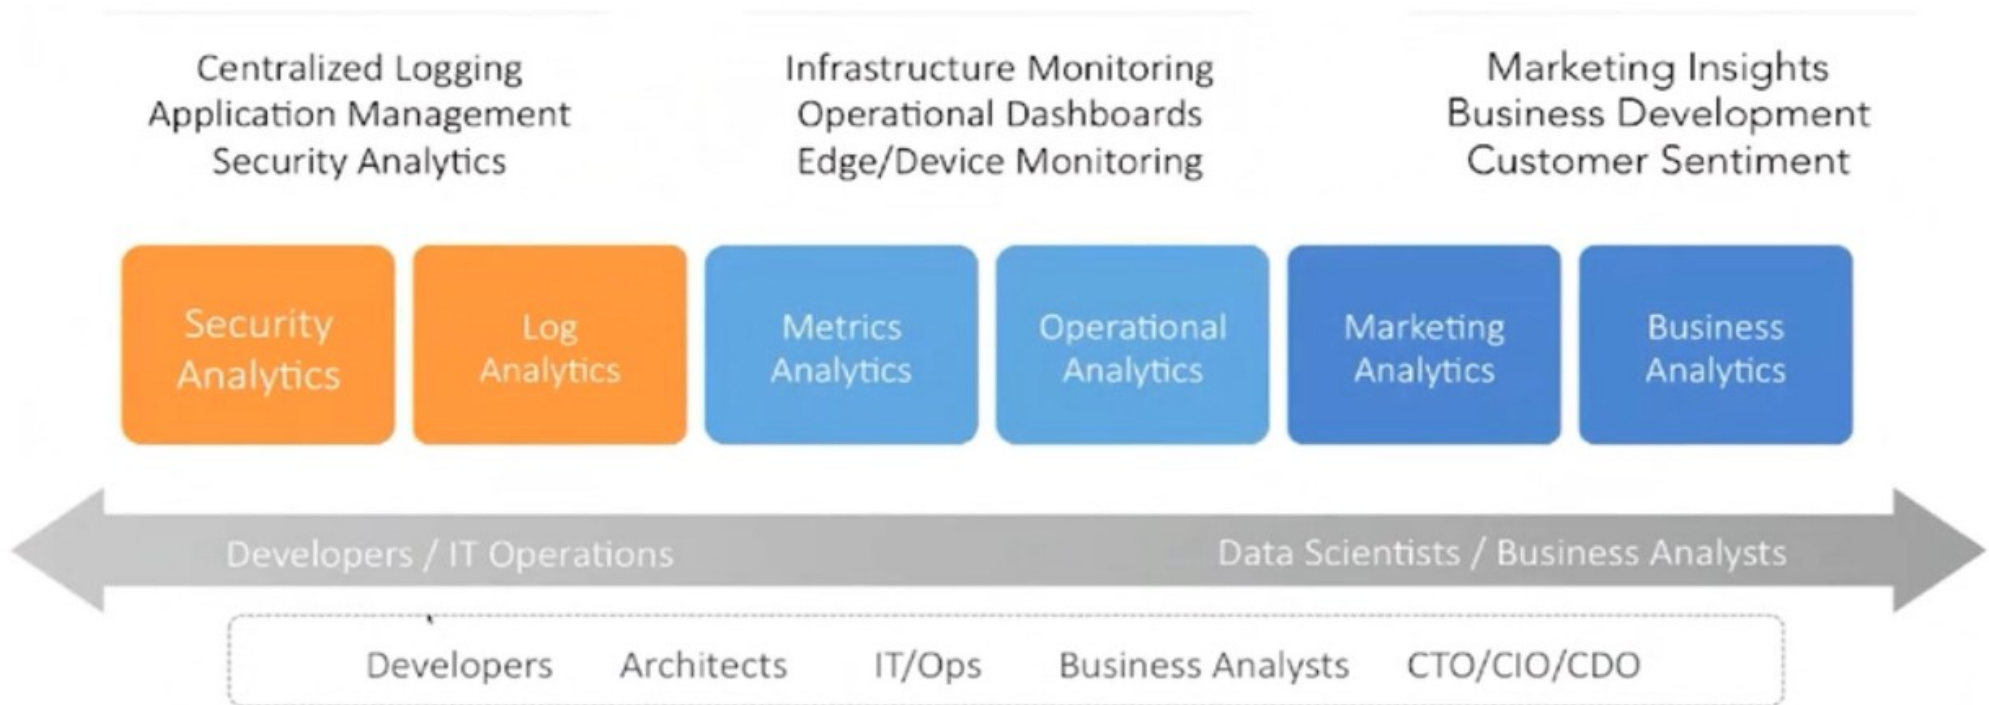
# Other Use Cases



The stack can have very different usages

- Business Analytics : Real-time analysis of KPI
- Marketing : User behavior, frequentation analysis, marketing segmentation
- Business risk management and fraud detection
- IoT (domotic or health sensors ...)
- Big Data in general

# Elastic Use cases





# Introduction

---

ELK proposition and use cases  
Distributions : Community vs X-Pack  
Built-in features  
**Deployment alternatives**  
Core concepts of clustering



# Sizing of the cluster

---

Depending on the use case, we can deploy different kind of cluster.

- Simple Search Engine with a reasonable amount of data :

A one node cluster may be sufficient

- Large amount of data and/or different indices  
A 3 nodes cluster and all the nodes have the same configuration
- Big Data :  
Cluster with different kind of specialized nodes (ingestion, queries, machine learning jobs, ...)
- Big Data and global company  
A cluster spread on different data centers



# Where to deploy ?

---

Standalone bare-metal or VMS operate by your own

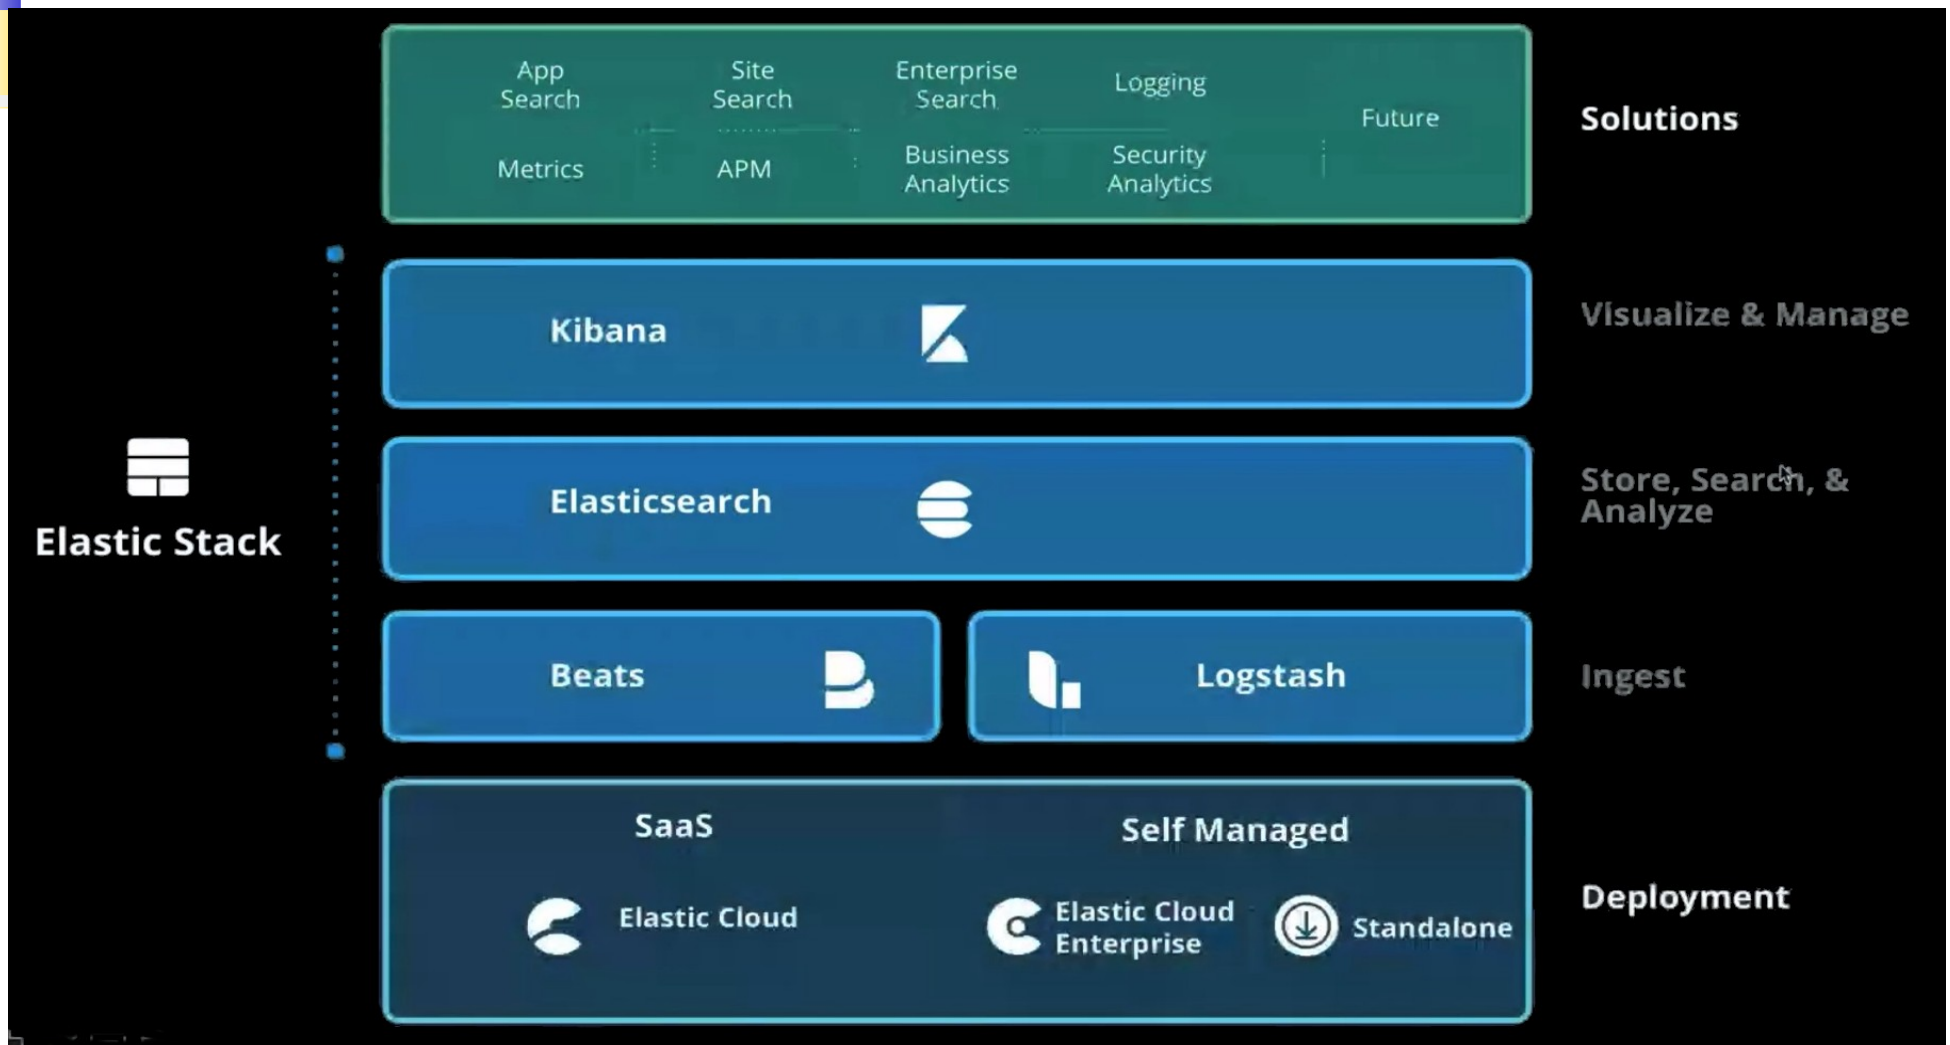
Private cloud (Kubernetes, OpenShift, OpenStack) operate by your own

Elastic Cloud provided by elastic company

Other public clouds also proposed ELK services



# In summary





# Introduction

---

ELK proposition and use cases  
Distributions : Community vs X-Pack  
Built-in features  
Deployment alternatives  
**Core concepts of clustering**



# Cluster

---

A **cluster** is a set of servers (nodes) that contains all the data and offers search capabilities on the different nodes

- It has a unique name in the local network (default : "*elasticsearch*").

=> A cluster can consist of only one node

=> A node cannot belong to 2 separate clusters



# Node

---

A node is a single server (Java process) that is part of a cluster.

A node generally stores data, and participates in the indexing and search functionalities of the cluster.

A node is also identified by a unique name (generated automatically if not specified)

The number of nodes in a cluster is not limited



# Master node

---

In a cluster, a node is elected as the **master node**, it is in charge of managing the configuration of the cluster and the creation of indexes.

For all document operations (indexing, searching), each data node is interchangeable and a client can address any of the nodes

In large architectures, nodes can be specialized for data ingestion, Machine Learning processes, and in this case, they do not participate in search functionalities.



# Index

---

An index is a collection of documents that have similar characteristics

- For example an index for customer data, another for the product catalog and yet another for orders

An index is identified by a name (in lower case)

- The name is used for indexing and search operations

In a cluster, you can define as many indexes as you want



# Document

---

A **document** is the basic unit of information that can be indexed.

- document has a set of fields (key/value) expressed with JSON format

Inside an index, you can store as many documents as you want



# Sample

---

```
{
  "name": "John Smith",      // String (may be analyzed or not)
  "age": 42,                 // Nombre
  "confirmed": true,         // Booléen
  "join_date": "2014-06-01", // Date
  "home": {                 // Imbrication
    "lat": 51.5,
    "lon": 0.1
  },
  "accounts": [             // tableau de données
    {
      "type": "facebook",
      "id": "johnsmith"
    }, {
      "type": "twitter",
      "id": "johnsmith"
    }
  ]
}
```





# Metadata

---

Metadata is also associated with each document.

Main metadata are :

- ***\_index*** : The location where the document is stored
- ***\_id*** : The unique identifier
- ***\_version*** : The version number of the document
- ...



# Shard

---

An index can store a very large amount of documents that may exceed the limits of a single node.

- To overcome this problem, ELS allows you to subdivide an index into several parts called **shards**

When creating the index, it is possible to define the number of shards

Each shard is an independent index that can be hosted on one of the cluster nodes



# Benefits of sharding

---

Sharding allows :

- To scale the volume of content
- To **distribute** and parallelize operations => improve performance



# Replica

---

To tolerate failures, it is recommended to use fail-over mechanisms in the event that a node fails (Hardware failures or upgrade of a system)

ELS allows you to set up copies of shards: **replicas** in order to

- **High-availability** : A replica is hosted on a different node than the primary shard
- To **scale** . Requests can be load balanced on every replica .



# Data Ingestion

---

## **Beats**

Logstash Concepts: pipelines, input,  
filters and output plugins  
Elastic search pipelines



# Introduction to Beats

---

Beats convey data

They are installed as agents on the different targets (servers) and periodically send their data

- Directly to *ElasticSearch*
- Or, if the data need to be transformed, to *Logstash*



# Types of beats

---

ELK provides :

- ***Packetbeat*** : All the network packets seen by a specific host
- ***Filebeat*** : All the lines written in a log file .
- ***Metricbeat*** : Performance metrics on CPU, Memory, Disk, IO, etc
- ***Winlogbeat*** : Windows log events

They all produce data which can be directly ingested in ElasticSearch

They are also distributed with Kibana dashboard which can easily be imported and used

Other beats provided by third party are also available



# Community beats

---

***amazonbeat*** : Reads data from a specified Amazon product.

***apachebeat, nginxbeat*** : Reads status from HTTPD server-status.

***browserbeat*** : Reads and ships browser history (Chrome, Firefox, & Safari) to an Elastic output.

***githubbeat*** : Easily monitors GitHub repository activity.

***httpbeat*** : Polls multiple HTTP(S) endpoints and sends the data to Logstash or Elasticsearch.

***kafkabeat*** : Reads data from Kafka topics.

***prometheusbeat*** : Send Prometheus metrics to Elasticsearch

***twitterbeat*** : Reads tweets for specified screen names.

...





# Data Ingestion

---

Beats

**Logstash Concepts: pipelines,  
input, filters and output plugins**

Elastic search pipelines



# Introduction

---

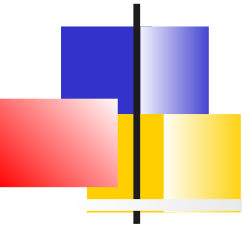
Logstash is a **real-time data collection** tool able to unify and normalize data from different sources

Originally focused on log files, it has evolved to handle very diverse events

Logstash is based on the notion of processing **pipelines**.

Extensible via plugins, it can read/write from/to very different sources and targets

# Sources of events



---

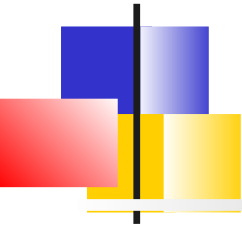
Traces and metrics : Server logs, syslog, Windows Event, JMX, ...

Web : HTTP request, Twitter, Hook for GitHub, JIRA, Polling d'endpoint HTTP

PersistenceStore : JDBC, NoSQL, \*MQ

Miscellaneous sensors : Mobile phones, Connected Home or Vehicles, Health sensors

# Data enrichment



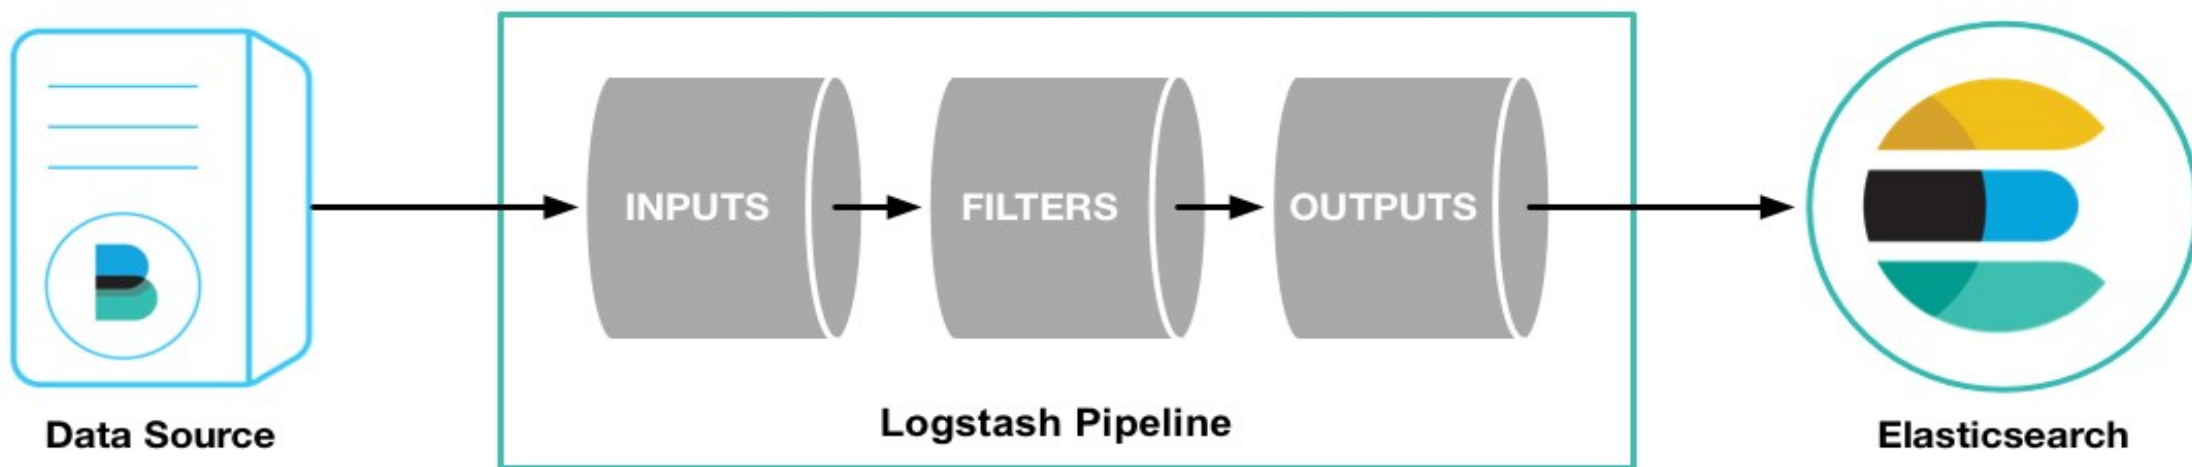
Logstash also allows you to enrich the input data.

- Geo-location from an IP address. (lookup to a service)
- Data encoding/decoding
- Date normalization
- Elastic Search queries for enrichment
- Anonymize sensitive information

# Pipeline Logstash

A logstash **pipeline** defines :

- An input plugin to read data from a source
- An output plugin to write data to a target (typically ElasticSearch)
- Optional filters to perform transformations





# Inputs

---

Some inputs :

- **beats** : Events produced by a beat
- **elasticsearch** : ElasticSearch Query response
- **exec** : Shell command output
- **file** : Reading lines of files
- **jdbc** : Records from database
- **http-poll, http, STOMP** : Events received via HTTP
- **Imap, Irc, rss** : Read mails, irc, rss, twitter
- **Kafka, redis, rabbitmq, jms** : Messages in message brokers
- **Tcp, udp, unix, syslog** : Low level socket events ...



# Some filters

---

***geoip*** : Ajoute des informations latitude/longitude à partir d'une IP

***grok*** : Divise un champ en d'autres champs via des expressions régulières

***mutate*** : Transforme un champ

***prune*** : Filtre des événements à partir d'une liste rouge ou blanche

***range*** : Vérifie que la taille ou longueur d'un champs est dans un intervalle

***translate*** : Remplace les valeurs des champs à partir d'une table de hash ou fichier YAML

***truncate*** : Tronque les champ supérieur à une certaine longueur

***urldecode*** : Décode les champs URL-encoded

***useragent*** : Parse les chaînes user agent

***xml*** : Parse le format XML



# Output plugins

---

## Some outputs

- **csv** : Write CSV files
- **elasticsearch** : Index into Elasticsearch
- **email** : Send the event via mail
- **file** : Writes into a file
- **stdout** : Writes in the standard output
- **pipe** : Send to standard input of another program
- **http** : Send to an HTTP or HTTPS endpoint
- **kafka, redis, rabbitMQ** : Messaging
- ...





# *Elasticsearch* Output

---

Main options of elasticsearch output are :

- **hosts** : Data nodes of the cluster
- **index** : The name of the index. Example : logstash-%  
{+YYYY.MM.dd}.
- **template** : A path to a template to create the index
- **template\_name** : The template's name present in Elasticsearch
- **document\_id** : The id of the Elasticsearch document (allow us updates)
- **pipeline** : The Elasticsearch pipeline to execute before indexing
- **routing** : To specify the shard to use



# Data Ingestion

---

Beats

Logstash Concepts: pipelines, input,  
filters and output plugins

**Elastic search pipelines**



# Introduction

---

Without logstash, it is possible to perform common transformations on the data before indexing by using the **ingest pipelines of Elasticsearch**

Features are similar to logstash but less powerful

ELS pipelines can be used to remove fields, extract values from text, and enrich your data.

It consists of a series of configurable tasks called **processors**. Each processor runs sequentially, making specific changes to incoming documents



# Pipeline

---

A **pipeline** is defined via :

- A name
- A description
- And a list of *processors*

Processors are predefined

They will be executed in sequence



# Usage

---

**#Create a pipeline named *attachment***

```
PUT _ingest/pipeline/attachment
{
  "description" : "Extract attachment information",
  "processors" : [
    { "attachment" : { "field" : "data" } }
  ]
}
```

**# Using a pipeline during indexation of a document.**

```
PUT my_index/my_type/my_id?pipeline=attachment
{
  "data":
  "e1xydGYxXGFuc2kNCkxvcmVtIGlwc3VtIGRvbG9yIHNpdCBhbWV0DQpccGFy
  IH0="
}
```



# Result

---

```
GET my_index/my_type/my_id
{
  "found": true,
  "_index": "my_index",
  "_type": "my_type",
  "_id": "my_id",
  "_version": 1,
  "_source": {
    "data":
    "e1xydGYxXGFuc2kNCkxvcmVtIGlwc3VtIGRvbG9yIHNpdCBhbWV0DQpccGFyIH0=",
    "attachment": {
      "content_type": "application/rtf",
      "language": "ro",
      "content": "Lorem ipsum dolor sit amet",
      "content_length": 28
    }
  }
}
```



# Available processors

---

Append : Add a field

Convert : Data type conversion

CSV : CSV format

Date : Create a @timestamp field

Drop : Drop the document on some condition

Enrich : Enrich data from other index

Fingerprint : Compute a fingerprint to detect doublon

GeoIP : Determine geo coordinates from IP Address

HTML strip : Remove HTML tags

Lowercase : Transform to lowercase some fields

Remove : Remove some fields

Rename : Rename a field

User agent : Split the user agent string



# ELS plugins

---

The functionalities of ELS can be increased via the notion of plugin

Plugins contain JAR files, but also scripts and configuration files.

They add new features, for example new pipeline's processor

They can be installed easily with

```
sudo bin/elasticsearch-plugin install [plugin_name]
```





# Ingest plugins

---

A particular ingest plugins allow to index office documents (Word, PDF, XLS, ...) :

- ***Attachment Plugin*** : Extracts text data from attachments in different formats (PPT, XLS, PDF, etc.). It uses the Apache Tika library.

```
sudo bin/elasticsearch-plugin install ingest-attachment
```



# Search Engine Capabilities

---

## **Inverted index and distributed search**

Supported Data Types

Analyzers, predined and custom

DSL syntax for queries

Full text and relevance score

Advanced search

Aggregations or faceting

Geo queries



# Elastic Search

---

2 techniques allows Elasticsearch to obtain amazing performance :

- The inverted index (Lucene)
- The distributed search among the nodes of the cluster



# Inverted index

---

In order to speed up searches on text fields, ELS uses a data structure called ***inverted index***

This consists of a unique word list where each word is associated with the documents in which it appears.

The method list is created after an analysis phase of the original text field.



# Example

---

	A	B
1	term	docs
2	pizza	3, 5
3	solr	2
4	lucene	2, 3
5	sourcesense	2, 4
6	paris	1, 10
7	tomorrow	1, 2, 4, 10
8	caffè	3, 5
9	big	6
10	brown	6
11	fox	6
12	jump	6
13	the	1, 2, 4, 5, 6, 8, 9



# Distributed search

---

Research requires a complex execution model, with documents scattered across shards

ElasticSearch must consult a copy of each shard of the index or indexes requested

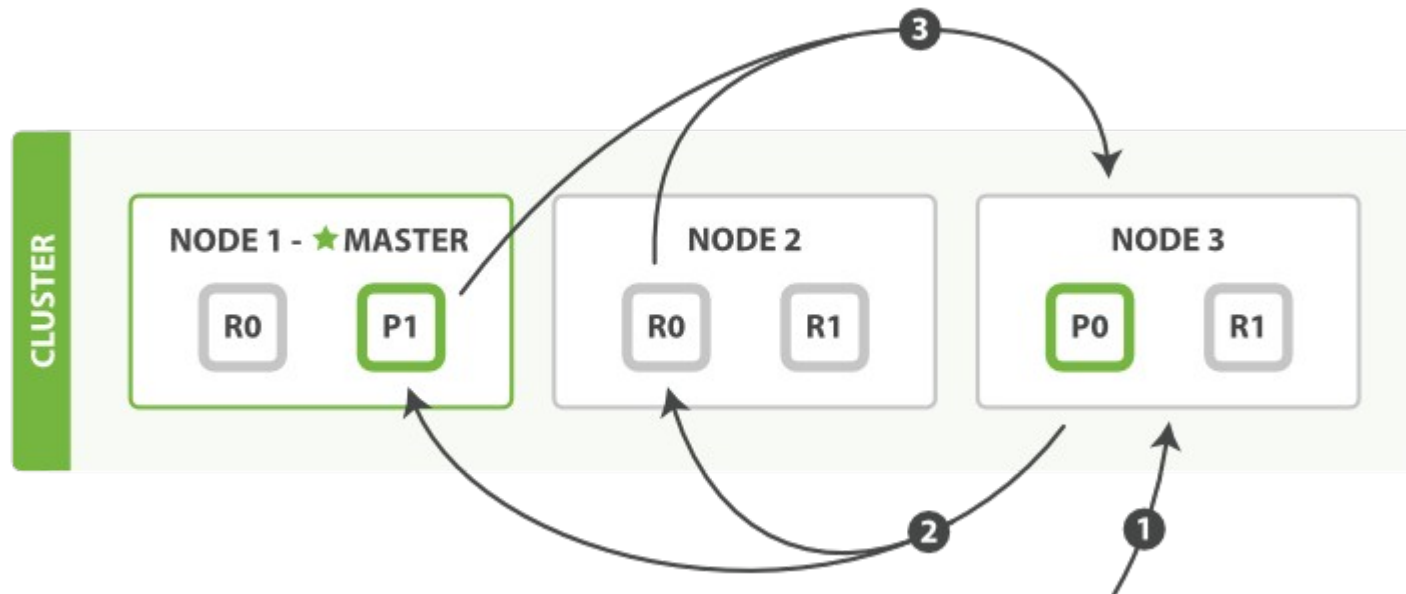
When found; results from different shards must be combined into a single list so that the API can return a page of results

The search is therefore executed in 2 phases:

- ***query***
- ***fetch.***

# Query phase

During the 1st phase, the request is broadcast to a copy (primary or replica) of all shards. Each shard performs the search and builds a priority queue of matching documents





# Query phase

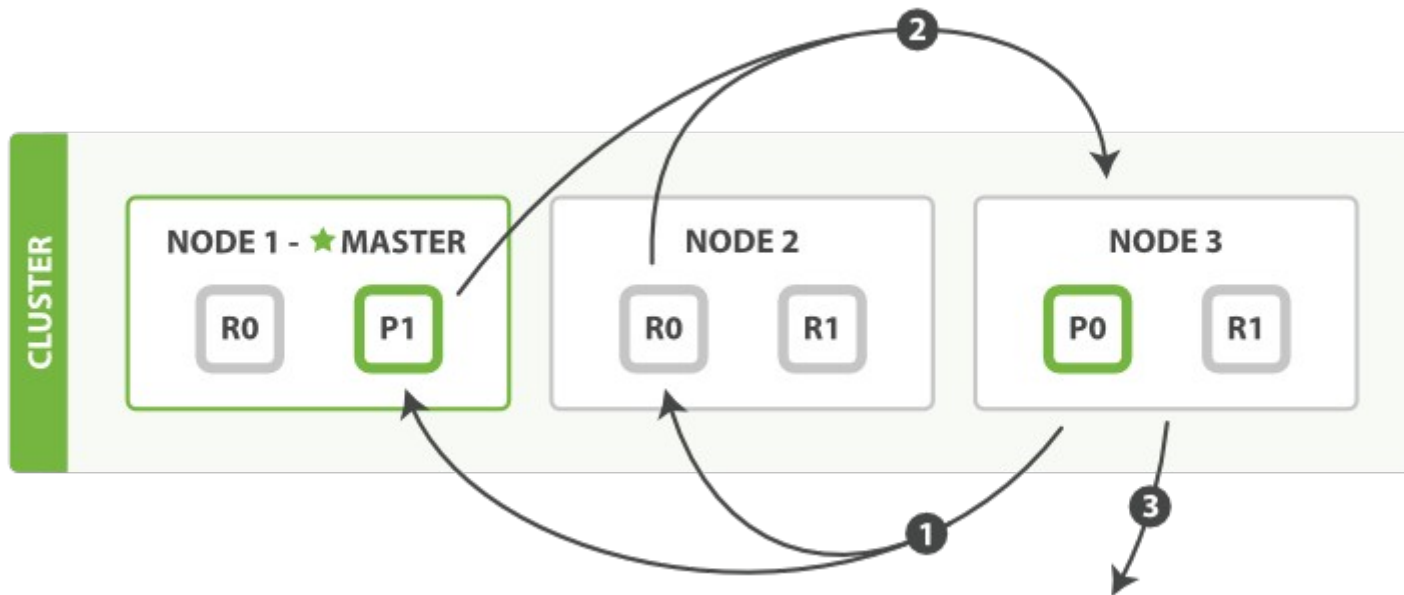
---

1. The client sends a query to node 3 which creates an empty priority queue of size = *from* + *size* .
2. Node 3 forwards the request to a copy of each shard in the index. Each shard executes the query locally and adds the result to a local priority queue of size = *from* + *size* .
3. Each shard returns the IDs and sort values of all documents in the queue to node 3 which merges these values into its priority queue



# Phase fetch

The fetch phase consists of retrieving the documents present in the priority queue.





# Fetch phase

---

1. The coordinator identifies which documents need to be retrieved and issues a multiple GET request to the shards.
2. Each shard loads the documents, enriches them if necessary (highlighting for example) and returns them to the coordinator node
3. When all documents have been retrieved, the coordinator returns the results to the client.



# Search Engine Capabilities

---

Inverted index and distributed search

## **Supported Data Types**

Analyzers, predined and custom

DSL syntax for queries

Full text and relevance score

Advanced search

Aggregations or faceting

Geo queries



# Introduction

---

Each field of a document has an associated datatype.

The metadata which specifies the different fields and their associated types is called the **mapping**

The mapping API is used to define, visualize, update the mapping of an index



# Supported simples data types

---

ELS supports the following simple data types :

- 2 kinds of string : *text* (analyzed) or *keyword* (not analyzed)
- Numbers : *byte* , *short* , *integer* , *long* , *float* , *double* , *token\_count*
- Booleans : *boolean*
- Dates : *date*
- Bytes : *binary*
- Range : *integer\_range* , *float\_range* , *long\_range* , *double\_range* , *date\_range*
- IP adress : *IPV4* ou *IPV6*
- Geolocation : *geo\_point* , *geo\_shape*
- A query (JSON structure) : *percolator*



# Complex data types

---

In addition to simple types, ELS supports

- **arrays** : There is no special mapping for arrays. Each field can contain 0, 1 or n values  

```
{ "tag": [ "search", "nosql" ] }
```

  - Array values must be of the same type
  - A null field is treated as an empty array
- **objects** : It is a data structure embedded in a field
  - Each sub-field is accessed with *object.property*



# Exact value or full-text

---

*String* stored by ELS can be of 2 data types :

- **keyword** : The value is taken as is, filter type operators can be used when searching  
Foo!= foo
- **text** : The value is analyzed and split into terms or tokens. Full-text search operators can be used when searching. This concerns data in natural language



# Mapping

---

ELS is able to dynamically generate the mapping during indexation

- It guesses the type of fields
- For string, it uses the value as text and keyword datatypes

Even if this feature is interesting for a proof or concept, in a real project you have to define the datatype of your index and to choose how the text string will be analyzed by the search Engine.

These decisions are important because you cannot change a mapping after indexing documents





# Search Engine Capabilities

---

Inverted index and distributed search

Supported Data Types

**Analyzers, predined and custom**

DSL syntax for queries

Full text and relevance score

Advanced search

Aggregations or faceting

Geo queries



# Tokenization and Normalization

---

To build the inverted index, ELS needs to separate a text into words (**tokenization**) and then **normalize** the words so that searching for the term "sending" for example returns documents containing: "sending", "Sending", "sendings", "send", ...

Tokenization and normalization are called **analysis**.

The analysis applies to the documents during indexing **AND** during the search on the search terms.



# Analysis steps

---

Analyzers transform a text into a stream of “tokens”. It is a combination of:

- **Character filters** prepare text by performing character replacement (& becomes *and*) or deletion (removal of HTML tags)
- **Tokenizer** . It splits a text into a series of lexical units: tokens (only 1 per analyzer)
- **Filters** take a token stream as input and transform it into another token stream



# Predefined analyzers

---

ELS offers directly usable analyzers:

- **Standard Analyzer** : This is the default analyzer. The best choice when the text is in various languages. It consists of :
  - Split text into words
  - Remove punctuation
  - Lowercase all words
  - Supports stop words
- **Simple analyzer** : Separates text based on character different than a letter, changes to lowercase
- **Stop analyzer** : Idem as simple but supports stop words
- **WhiteSpace Analyzer**: Separates text based on spaces
- **Pattern analyzer** : Based on regular expression
- **Language analyzer**: These are language-specific parsers. They include "stop words" (remove most common words) and extract the root of a word. This is the best choice if the index is in one language
- **Fingerprint analyzer** : Use for duplicate detection



# Testing analyzers

---

Every analyzers can be tested via the API :

GET /\_analyze?analyzer=standard  
Text to analyze

Réponse :

```
{
  "tokens": [ {
    "token": "text",
    "start_offset": 0,
    "end_offset": 4,
    "type": "<ALPHANUM>",
    "position": 1
  }, {
    "token": "to",
    "start_offset": 5,
    "end_offset": 7,
    "type": "<ALPHANUM>",
    "position": 2
  }, {
    "token": "analyze",
    "start_offset": 8,
    "end_offset": 15,
    "type": "<ALPHANUM>",
    "position": 3
  } ] }
```



# Specify mapping at index creation

---

```
PUT /office2
{
  "mappings": {
    "dynamic": false,
    "properties": {
      "author": { "type": "text" },
      "content": {
        "type": "text",
        "analyzer": "english",
        "fields": {
          "fr": {
            "type": "text",
            "analyzer": "french"
          }
        }
      },
      "content_length": { "type": "long" },
      "date": { "type": "date" },
      "description": { "type": "text" },
      "language": { "type": "keyword" },
      "title": { "type": "text" }
    }
  }
}
```

2 fields are available for full-text searching : *content* and *content.fr*



# Custom analyzers

---

It is possible to define yours custom analyzers

- Either by defining precisely the tokenizer, character filters and filters to use
- Either by overriding an existing analyzer



# Example

---

```
PUT /my_index
{
  "settings": {
    "analysis": {
      "char_filter": {
        "&_to_and": {
          "type": "mapping",
          "mappings": [ "&=> and " ]
        },
      },
      "filter": {
        "my_stopwords": {
          "type": "stop",
          "stopwords": [ "the", "a" ]
        },
      },
      "analyzer": {
        "my_analyzer": {
          "type": "custom",
          "char_filter": [ "html_strip", "&_to_and" ],
          "tokenizer": "standard",
          "filter": [ "lowercase", "my_stopwords" ]
        }
      }
    }
  }
}
```





# Predefined filters

---

ELS provides many filters for building custom analyzers. For example :

- **Length Token** : Deleting words that are too short or too long
- **N-Gram** and **Edge N-Gram** : Analyzer to speed up search suggestions
- **Stemming filters** : Algorithm to extract the root of a word
- **Phonetic filters** : Phonetic representation of words
- **Synonym** : Word Match
- **Keep Word** : The opposite of stop words
- **Limit Token Count** : Limit the number of tokens associated with a document
- **Elison Token** : Handling of apostrophes (exemple : French)



# Synonym filter

---

Synonyms can be used to merge words that have nearly the same meaning.

Ex : pretty, cute, beautiful

They can also be used to make a word more generic.

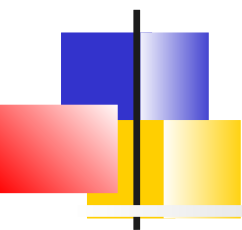
- For example, bird can be used as a synonym for pigeon, sparrow, ...



# Example

---

```
PUT /my_index
{
  "settings": {
    "analysis": {
      "filter": {
        "my_synonym_filter": {
          "type": "synonym",
          "synonyms": ["british,english", "queen,monarch"]
        }
      },
      "analyzer": {
        "my_synonyms": { "tokenizer": "standard",
                        "filter": ["lowercase","my_synonym_filter"]
        }
      }
    }
  }
}
```



# Search Engine Capabilities

---

Inverted index and distributed search

Supported Data Types

Analyzers, predined and custom

**DSL syntax for queries**

Full text and relevance score

Advanced search

Aggregations or faceting

Geo queries



# Introduction

---

Searches with a request body offer more functionalities than the search lite mechanism.

In particular, they allow you to :

- Combine query clauses more easily
- To influence the score
- Highlight parts of the result
- Aggregate subsets of results
- Return suggestions to the user



# Combination example

---

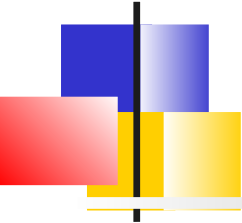
```
{  
  "bool": {  
    "must": { "match": { "tweet": "elasticsearch" } },  
    "must_not": { "match": { "name": "mary" } },  
    "should": { "match": { "tweet": "full text" } }  
  }  
}
```



# Ex: Combination of combinations

---

```
{
  "bool": {
    "must": {
      "match": { "email": "business opportunity" }
    },
    "should": [
      { "match": { "starred": true } },
      { "bool": {
        "must": { "folder": "inbox" } },
        "must_not": { "spam": true } }
      ]
    },
    "minimum_should_match": 1
  }
}
```



# Distinction between query and filter

---

## Distinction between query and filter

- **Filters** are used for fields with exact values.  
Their result is of Boolean type, i.e. a document satisfies a filter or not  
They are very efficient
- **Queries** calculate a relevance score for each document found.  
The result is sorted by the relevance score by default  
More computation but more relevant





# Types of operators

---

Operators can be classified in 2 families :

- Operators that do not perform analysis and operate on a single term (*term*, *fuzzy*).
- The operators that apply the analyzer to the search terms corresponding to the searched field (*match*, *query\_string*) .



# Operators without analysis

---

***term*** : Used to filter exact values :

```
{ "term": { "age": 26 } }
```

***terms*** : Allows multiple values to be specified :

```
{ "terms": { "tag": [ "search", "full_text",  
"nosql" ] } }
```

***range*** : Allows you to specify a date or numeric range:

```
{ "range": { "age": { "gte": 20, "lt":  
30 } } }
```

***exists*** and ***missing*** : Allows you to test whether or not a document contains a field

```
{ "exists": { "field": "title" } }
```



# *match*

---

The ***match*** operator is the standard operator to execute full-text search.

```
{ "match": { "tweet": "About Search" } }
```

It is a Boolean query which by default analyzes the words passed as parameters and builds an OR type query.



# *multi\_match*

---

The ***multi\_match*** operator allows you to run the same query on multiple fields:

```
{"multi_match": { "query": "full text search", "fields":  
[ "title", "body" ] } }
```

Wildcards can be used for fields

Fields can be boosted with notation ^

```
GET /_search  
{  
  "query": {  
    "multi_match" : {  
      "query" : "this is a test",  
      "fields" : [ "subject^3", "text*" ]  
    }  
  }  
}
```



# *query\_string*

---

***query\_string*** operator uses a parser to understand the query string. (the same as for lite search)

It has many parameters (default\_field, analyzer, ...)

```
GET /_search
{
  "query": {
    "query_string" : {
      "default_field" : "content",
      "query" : "title:(quick OR brown)"
    }
  }
}
```



## *simple\_query\_string*

---

The ***simple\_query\_string*** operator avoids causing a parsing error in the search query.

It can be used directly with what a user has entered in a search field

It understands the simplified syntax:

- + for AND
- | for OR
- - for negation
- ...



# Search Engine Capabilities

---

Inverted index and distributed search

Supported Data Types

Analyzers, predined and custom

DSL syntax for queries

**Full text and relevance score**

Advanced search

Aggregations or faceting

Geo queries



# Sorting

---

During filter-type queries, it may be interesting to set the sorting criterion.

This is done via the **sort** parameter

GET /\_search

```
{
  "query" : {
    "bool" : {
      "filter" : { "term" : { "user_id" : 1 } }
    }
  }, "sort": { "date": { "order": "desc" } }
}
```





# Response

In the response, the `_score` field is not calculated and the value of the `sort` field is specified for each document

```
"hits" : {
  "total" : 6,
  "max_score" : null,
  "hits" : [ {
    "_index" : "us",
    "_type" : "tweet",
    "_id" : "14",
    "_score" : null,
    "_source" : {
      "date": "2014-09-24",
      ...
    }, "sort" : [ 1411516800000 ]
  },
  ...
}
```



# Other sorting features

---

It is possible to combine a sorting criterion with another criterion or the score. Please note the order is important.

It is also possible to use an aggregate function to sort multivalued fields (*min*, *max*, *sum*, *avg*, ...)



# Relevance

---

The score of each document is represented by a floating point number (`_score`). The higher this number, the more relevant is the document.

The algorithm used by ELS is called ***term frequency/inverse document frequency***, or ***TF/IDF***.

It takes into account the following factors:

- The ***frequency of the term*** : how many times the term appears in the field
- The ***inverse document frequency*** : How many times does the term appear in the index? The more the term appears, the less weight it has.
- The ***normalized length of the field*** : The longer the field, the less relevant the terms

Depending on the type of query (fuzzy query, boost factor ...) other factors can influence the score



# Explanation of relevance

---

With the ***explain*** parameter, ELS provides an explanation of the score

```
GET /_search?explain
```

```
{ "query"
: { "match" : { "tweet" : "honeymoon" }}
}
```

The *explain* parameter can also be used to understand why a document matches or not.

```
GET /us/12/_explain
```

```
{
"query" : {
  "bool" : {
    "filter" : { "term" : { "user_id" : 2 }},
    "must" : { "match" : { "tweet" : "honeymoon" }}
  } } }
```



# Response

---

```
"_explanation": {
  "description": "weight(tweet:honeyymoon in 0)
[PerFieldSimilarity], result of:",
  "value": 0.076713204,
  "details": [ {
    "description": "fieldWeight in 0, product of:",
    "value": 0.076713204,
    "details": [ {
      "description": "tf(freq=1.0), with freq of:",
      "value": 1,
      "details": [ {
        "description": "termFreq=1.0",
        "value": 1
      } ]
    }, {
      "description": "idf(docFreq=1, maxDocs=1)",
      "value": 0.30685282
    }, {
      "description": "fieldNorm(doc=0)",
      "value": 0.25,
    } ]
  } ]
} ] }
```



# *multi\_match* and relevance score

---

With the operator *multi\_match*, the ***type*** parameter specifies the way the score is calculated.

- ***best\_fields*** (default) : Find documents that match one of the fields but use the best field to assign the score
- ***most\_fields*** : Combines the score of each field
- ***cross\_fields*** : Concatenates all fields and uses the same parser
- ***phrase*** : Use a *match\_phrase* operator on each field and combine each field's score
- ***phrase\_prefix*** : Use a *match\_phrase\_prefix* operator on each field and combine each field's score



# *bool* and relevance score

---

```
GET /my_index/my_type/_search
{ "query": {
  "bool": {
    "must": { "match": { "title": "quick" }},
    "must_not": { "match": { "title": "lazy" }},
    "should": [
      { "match": { "title": "brown" }},
      { "match": { "title": "dog" }}
    ]
  }
}
```

The calculation of relevance is done by adding the score of each *must* or *should* clause and dividing by 3



# Boosting clause

---

It is possible to give more weight to a particular clause by using the **boost** parameter

```
GET /_search
{
  "query": { "bool": {
    "must": { "match": {
      "content": { "query": "full text search", "operator": "and" }}}},
    "should": { "match": {
      "content": { "query": "Elasticsearch", "boost": 3 }}}
  } } }
```





# *boosting* operator

---

The ***boosting*** operator allows you to specify a clause that reduces the score of matching documents

```
GET /_search
```

```
{
  "query": {
    "boosting" : {
      "positive" : {
        "term" : {"text" : "apple"}
      },
      "negative" : {
        "term" : { "text" : "pie tart fruit crumble tree" }
      },
      "negative_boost" : 0.5 // requis le malus au document qui matche
    }
  }
}
```



# Function score query

---

The `function_score` allows you to modify the score of documents that are retrieved by a query.

Different functions to evaluate the score can be used :

- **`script_score`** : customize the scoring of with a computation derived from other numeric field values in the doc using a script expression
- **`weight`** : allows you to multiply the score by the provided weight
- **`random_score`** generates scores that are uniformly distributed
- **`field_value_factor`** : function allows you to use a field from a document to influence the score. It's similar to using the *`script_score`*



# Example : script\_score

---

```
GET /_search
{
  "query": {
    "function_score": {
      "query": {
        "match": { "message": "elasticsearch" }
      },
      "script_score": {
        "script": {
          "source": "Math.log(2 + doc['views'].value)"
        }
      }
    }
  }
}
```



# Example *Field Value factor*

---

```
GET /_search
{
  "query": {
    "function_score": {
      "field_value_factor": {
        "field": "my-int",
        "factor": 1.2,
        "modifier": "sqrt",
        "missing": 1
      }
    }
  }
}
```



# Search Engine Capabilities

---

Inverted index and distributed search

Supported Data Types

Analyzers, predined and custom

DSL syntax for queries

Full text and relevance score

**Advanced search**

Aggregations or faceting

Geo queries



# Partial Matching

---

Partial matching allows users to specify a portion of the term they are looking for

Common use cases are:

- Matching postal codes, serial numbers or other not\_analyzed values that start with a particular prefix or even a regular expression
- On-the-fly search: searches performed on each character typed to make suggestions to the user
- Matching in languages like German that contain long compound nouns



# *prefix* filter

---

The *prefix* filter is a search performed on the term. It does not parse the search string and assumes that the correct prefix has been provided.

```
GET /my_index/address/_search
{
  "query": {
    "prefix": { "postcode": "W1" }
  }
}
```



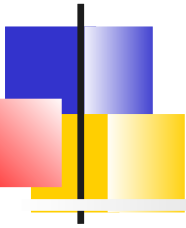
# Wildcard and regexp

**wildcard** or **regexp** searches are similar to *prefix* but allow the use of wildcards or regular expressions

```
GET /my_index/address/_search
{
  "query": {
    "wildcard": { "postcode": "W?F*HW" }
  }
}
```

```
GET /my_index/address/_search
{
  "query": {
    "regexp": { "postcode": "W[0-9].+" }
  }
}
```





# Indexing for auto-completion

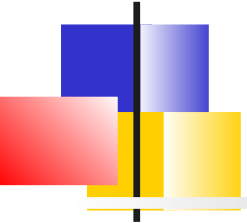
---

The principle is to index the beginnings of words of each term

This can be done through a particular ***edge\_ngram*** filter

```
{  
  "filter": {  
    "autocomplete_filter": { "type": "edge_ngram",  
                             "min_gram": 1,  
                             "max_gram": 20  
    }  
  }  
}
```

For each term, it creates n tokens of minimum size 1 and maximum 20. The tokens are the different prefixes of the term.



# *match\_phrase*

---

The ***match\_phrase*** operator parses the search string to produce a list of terms but only keeps documents that contain all of the terms in the same position.

```
GET /my_index/my_type/_search
{
  "query": {
    "match_phrase": { "title": "quick brown
fox" }
  }
}
```



# Proximity and *slop* parameter

---

It is possible to introduce flexibility to phrase matching by using the *slop* parameter which indicates how far apart the terms can be.

Documents with these closest terms will have better relevance.

```
GET /my_index/my_type/_search
{
  "query": {
    "match_phrase": {
      "title": {
        "query": "quick fox",
        "slop": 20 // ~ distance in words
      }
    }
  }
}
```



# *match\_phrase\_prefix*

The ***match\_phrase\_prefix*** operator behaves like *match\_phrase*, except it treats the last word as a prefix

It is possible to limit the number of expansions by setting the *max\_expansions* parameter

```
{  
  "match_phrase_prefix" : {  
    "brand" : {  
      "query": "johnnie walker bl",  
      "max_expansions": 50  
    }  
  }  
}
```



# Fuzzy query

---

**Fuzzy query** returns documents that contain terms similar to the search term, as measured by a Levenshtein edit distance.

An edit distance is the number of one-character changes needed to turn one term into another. These changes can include:

- Changing a character (box → fox)
- Removing a character (black → lack)
- Inserting a character (sic → sick)
- Transposing two adjacent characters (act → cat)

Most typos or misspelling errors have a distance of 1.

=> So 80% of errors could be fixed with single character editing



# Fuzzy query

---

It is possible during a request to position the **fuzziness** parameter which therefore gives the Levenshtein distance tolerance.

This parameter can also be set to AUTO:

- 0 for strings of 1 or 2 characters
- 1 for strings of 3, 4 or 5 characters
- 2 for strings longer than 5 characters

Be careful: No stemmer with fuzzy queries!



# Example

---

```
GET /my_index/my_type/_search
```

```
{  
  "query": {  
    "multi_match": {  
      "fields": [ "text", "title" ],  
      "query": "SURPRIZE ME!",  
      "fuzziness": "AUTO"  
    }  
  }  
}
```



# Highlighting

---

**Highlighting** consists of highlighting the matched term in the original content that has been indexed.

This is possible thanks to the metadata stored during indexing

The original field must be stored by ELS





# Example

---

```
GET /_search
{
  "query" : {
    "bool" : {
      "must" : { "match" :
{ "attachment.content" : "Administration" } },
      "should" : { "match" : { "attachment.content" :
"Oracle" } }
    }
  },
  "highlight" : {
    "fields" : {
      "attachment.content" : {}
    }
  }
}
```



# Response

---

```
"highlight" : {  
    "attachment.content" : [  
        " formation <em>Administration</em> <em>Oracle</em>  
Forms/Report permet de voir tous les aspects nécessaires à",  
        " aux autres services Web de l'offre  
<em>Oracle</em> Fusion Middleware. Cependant, l'administration  
de ces",  
        " Surveillance. Elle peut-être un complément de la  
formation « <em>Administration</em> d'un serveur Weblogic",  
        "-forme <em>Oracle</em> Forms ou des personnes  
désirant migrer leur application client serveur Forms vers",  
        " <em>Oracle</em> Forms 11g/12c\nPré-requis : \  
nExpérience de l'administration système\nLa formation  
« Administation"  
    ]  
}
```



# Tags used in the response

---

```
GET /_search
{
  "query" : {
    "bool" : {
      "must" : { "match" :
{ "attachment.content" : "Administration" } },
      "should" : { "match" : { "attachment.content" : "Oracle" } }
    }
  },
  "highlight" : {
    "pre_tags" : ["<tag1>"],
    "post_tags" : ["</tag1>"],
    "fields" : {
      "attachment.content" : { }
    }
  }
}
```



# Fragments

---

It is possible to control the highlighted fragments for each field:

- ***fragment\_size*** : gives the max size of the highlighted fragment
- ***number\_of\_fragments*** : The maximum number of fragments in this field



# Search Engine Capabilities

---

Inverted index and distributed search

Supported Data Types

Analyzers, predined and custom

DSL syntax for queries

Full text and relevance score

Advanced search

**Aggregations or faceting**

Geo queries



# Introduction

---

Aggregations are extremely powerful for reporting and dashboards

Using the Elastic, Logstash, and Kibana stack demonstrates just how much you can do with aggregations.

# Kibana dashboard





# DSL Syntax

---

An aggregation can be seen as a unit of work that builds analytical information on a set of documents.

Depending on its position in the DSL tree, it applies to all search results or to subsets

In DSL syntax, an aggregation block uses the keyword ***aggs***

**// The max of the price field in all documents**

```
POST /sales/_search?size=0
```

```
{
  "aggs" : {
    "max_price" : { "max" : { "field" : "price" } }
  }
}
```



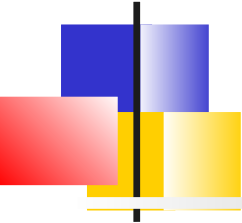


# Types of aggregations

---

Several concepts relate to aggregations:

- **Groups or Buckets** : Set of documents that have a field with the same value or share the same criteria.  
Groups can be nested. ELS provides syntaxes for defining groups and counting the number of documents in each category
- **Metrics**: Metric calculations on a group of documents (min, max, avg, ..)
- **Pipeline** : Aggregations that take input from other aggregations instead of documents or fields.



# Example Bucket

---

GET /cars/transactions/\_search

```
{  
  "aggs" : {  
    "colors" : {  
      "terms" : { "field" : "color.keyword" }  
    } } }  
}
```



# Response

---

```
{  
  ...  
  "hits": { "hits": [] },  
  "aggregations": {  
    "colors": {  
      "doc_count_error_upper_bound": 0, // incertitude  
      "sum_other_doc_count": 0,  
      "buckets": [  
        { "key": "red", "doc_count": 4 },  
        { "key": "blue", "doc_count": 2 },  
        { "key": "green", "doc_count": 2 }  
      ]  
    }  
  }  
}
```



# Example metric

---

GET /cars/transactions/\_search

```
{
  "size": 0,
  "aggs" : {
    "avg_price" : {
      "avg" : {"field" : "price"}
    }
  }
}
```



# Response

---

```
"hits": {  
  "total": 8,  
  "max_score": 0,  
  "hits": []  
},  
"aggregations": {  
  "avg_price": {  
    "value": 26500  
  }  
}
```



# Bucket/Metrics juxtaposition

---

```
GET /cars/transactions/_search
```

```
{
  "size": 0,
  "aggs" : {
    "colors" : {
      "terms" : { "field" : "color.keyword" }
    },
    "avg_price" : {
      "avg" : {"field" : "price"}
    }
  }
}
```



# Response

---

```
{  
  ...  
  "hits": { "hits": [] },  
  "aggregations": {  
    "avg_price": {  
      "value": 26500  
    },  
    "colors": {  
      "doc_count_error_upper_bound": 0,  
      "sum_other_doc_count": 0,  
      "buckets": [  
        { "key": "red", "doc_count": 4 },  
        { "key": "blue", "doc_count": 2 },  
        { "key": "green", "doc_count": 2 }  
      ]  
    }  
  }  
}
```



# Nesting aggregations

---

```
GET /cars/transactions/_search {  
  "aggs": {  
    "colors": {  
      "terms": { "field": "color" },  
      "aggs": {  
        "avg_price": {  
          "avg": { "field": "price" }  
        }, "make": {  
          "terms": { "field": "make" }  
        }  
      }  
    }  
  }  
}
```





# Response

---

```
{
  ...
  "aggregations": {
    "colors": {
      "buckets": [
        { "key": "red",
          "doc_count": 4,
          "avg_price": {
            "value": 32500
          },
          "make": { "buckets": [
            { "key": "honda", "doc_count": 3 },
            { "key": "bmw", "doc_count": 1 }
          ]
        }
      ],
    },
    ...
  }
}
```



# Aggregation and search

---

In general, an aggregation is combined with a search. The buckets are then deduced from the documents that match.

```
GET /cars/transactions/_search
{
  "query" : {
    "match" : { "make" : "ford" }
  }, "aggs" : {
    "colors" : {
      "terms" : { "field" : "color" }
    }
  }
}
```



# Sorting buckets

---

```
GET /cars/transactions/_search
```

```
{
```

```
  "aggs" : {
```

```
    "colors" : {
```

```
      "terms" : { "field" : "color",
```

```
                  "order": { "avg_price" : "asc" }
```

```
    }, "aggs": {
```

```
      "avg_price": {
```

```
        "avg": {"field": "price"}
```

```
      }
```

```
    } } } }
```



# Pipeline

---

Pipeline aggregations work on the outputs produced from other aggregations

They reference the aggregations used to perform their computation by using the ***buckets\_path*** parameter



# Example

---

POST /\_search

```
{
  "aggs": {
    "my_date_histo": {
      "date_histogram": {
        "field": "timestamp",
        "calendar_interval": "day"
      },
      "aggs": {
        "the_sum": {
          "sum": { "field": "lemmings" }
        },
        "the_deriv": {
          "derivative": { "buckets_path": "the_sum" }
        }
      }
    }
  }
}
```



# Types of buckets

---

ELS offers different ways of grouping data :

- By term: Requires field tokenization
- Histogram : By range of values
- Date Histogram : By date range
- By IP range
- By absence/presence of a field
- Significant terms
- By geo-location
- ..



# Histogram

---

GET /cars/transactions/\_search

```
{
  "aggs": {
    "price": {
      "histogram": {
        "field": "price",
        "interval": 20000
      },
      "aggs": {
        "revenue": {
          "sum": { "field" : "price" }
        }
      }
    }
  }
}
```



# Date histogram

---

```
GET /cars/transactions/_search
```

```
{
```

```
  "aggs": {
```

```
    "sales": {
```

```
      "date_histogram": {
```

```
        "field": "sold",
```

```
        "interval": "month",
```

```
        "format": "yyyy-MM-dd"
```

```
      }
```

```
    } } }
```





# *significant\_terms*

---

The ***significant\_terms*** aggregation is more subtle but can give interesting results (anomaly detection).

This consists of analyzing the returned data and finding the terms that appear at an abnormally higher frequency

Abnormally means: relative to the frequency for all documents

=> These statistical anomalies generally reveal interesting things

# Mechanism



*significant\_terms* takes a search result and performs an aggregation

It then starts from the set of documents and performs the same aggregation

It then compares the results of the first search that are "abnormal" against the overall search

With this type of operation, you can:

- People who liked... also liked...
- Customers who had questionable credit card transactions all went to such and such a merchant
- Every Thursday evening, the page is much more consulted



# Example

---

```
{  
  "query" : {  
    "terms" : { "force" : [ "British Transport Police" ] }  
  },  
  "aggregations" : {  
    "significantCrimeTypes" : {  
      "significant_terms" : { "field" : "crime_type" }  
    }  
  }  
}
```



# Response

---

```
"aggregations" : {  
  "significantCrimeTypes" : {  
    "doc_count": 47347, // Total query result  
    "buckets" : [  
      {  
        "key": "Bicycle theft",  
        "doc_count": 3640, // Number docs for the query result  
        "score": 0.371235374214817,  
        "bg_count": 66799 // Number for all documents  
      }  
      ...  
    ]  
  }  
}
```

=> Bike theft rate unusually high for « British Transport Police »

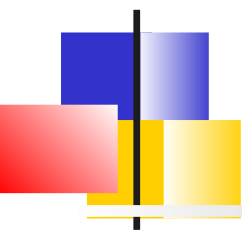


# Available metrics

---

ELS offers many metrics:

- ***avg, min, max, sum***
- ***value\_count, cardinality*** : Distinct value count
- ***top\_hit*** : The top documents
- ***extended\_stats*** : Statistical metrics (count, sum, variance, ...)
- ***percentiles*** : percentiles



# Search Engine Capabilities

---

Inverted index and distributed search

Supported Data Types

Analyzers, predined and custom

DSL syntax for queries

Full text and relevance score

Advanced search

Aggregations or faceting

**Geo queries**



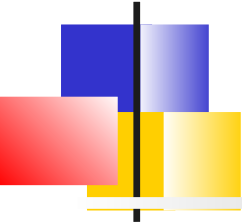
# Introduction

---

ELS allows you to combine geo-location with full-text, structured searches and aggregations

ELS has 2 data types to represent geolocation data

- ***geo\_point*** : represents a latitude-longitude couple. This mainly allows the calculation of distance
- ***geo\_shape*** : defines an area. This allows you to know if 2 areas have an intersection



# Geo-point

Geo-points cannot be automatically detected by ELS. They must be explicitly specified in the mapping:

```
PUT /attractions
```

```
{ "mappings": {  
  "properties": {  
    "name": { "type": "string" },  
    "location": { "type": "geo_point" }  
  }  
}
```

```
----
```

```
PUT /attractions/1
```

```
{"name": "Chipotle Mexican Grill", "location": "40.715, -74.011" }
```

```
PUT /attractions/2
```

```
{ "name": "Pala Pizza", "location": { "lat":40.722,"lon": -73.989 } }
```

```
PUT /attractions/3
```

```
{ "name": "Mini Munchies Pizza", "location": [ -73.983, 40.719 ] }
```





# Filters

---

4 filters can be used to include or exclude documents with respect to their *geo-point*:

- ***geo\_bounding\_box***: The geo-points included in the provided rectangle
- ***geo\_distance***: Distance from a center point below a boundary.  
Sorting and scoring can be relative to distance
- ***geo\_distance\_range***: Distance in a range
- ***geo\_polygon***: Geo-points include in a polygon



# Example

---

GET /attractions/restaurant/\_search

```
{
  "query": {
    "bool": {
      "filter": {
        "geo_bounding_box": {
          "location": { "top_left": { "lat": 40.8, "lon": -74.0 },
                        "bottom_right": { "lat": 40.7, "lon": -73.0 }
          }
        }
      }
    }
  }
}
```



# Aggregation

---

3 types of aggregation on geo-points are possible

- ***geo\_distance*** (bucket): Groups documents in concentric circles around a central point
- ***geohash\_grid*** (bucket): Group documents by cells (*geohash\_cell*, google maps squares) for display on a map
- ***geo\_bounds*** (metrics): returns the coordinates of a rectangle area that would encompass all geo-points. Useful for choosing the right zoom level



# Example

---

```
GET /attractions/restaurant/_search
{
  "query": { "bool": { "must": {
    "match": { "name": "pizza" }
  },
  "filter": { "geo_bounding_box": {
    "location": { "top_left": { "lat": 40.8, "lon": -74.1 },
                  "bottom_right": { "lat": 40.4, "lon": -73.7 }
    }
  } } } },
  "aggs": {
    "per_ring": {
      "geo_distance": {
        "field": "location",
        "unit": "km",
        "origin": {
          "lat": 40.712,
          "lon": -73.988
        },
        "ranges": [
          { "from": 0, "to": 1 },
          { "from": 1, "to": 2 }
        ]
      }
    }
  }
}
```



# Geo-shape

---

Like *geo\_point* , ***geo-shape*** fields must be mapped explicitly:

```
PUT /attractions
```

```
{  
  "mappings": { "landmark": {  
    "properties": {  
      "name": { "type": "string" },  
      "location": { "type": "geo_shape" }  
    } } } }  
}
```

```
---
```

```
PUT /attractions/landmark/dam_square
```

```
{  
  "name" : "Dam Square, Amsterdam",  
  "location" : {  
    "type" : "polygon",  
    "coordinates" : [[ [ 4.89218, 52.37356 ], [ 4.89205, 52.37276 ], [ 4.89301, 52.37274 ],  
      [ 4.89392, 52.37250 ], [ 4.89218, 52.37356 ] ] ]  
  } }  
}
```



# Query example

---

```
GET /attractions/landmark/_search
{
  "query": {
    "geo_shape": {
      "location": {
        "shape": {
          "type": "circle",
          "radius": "1km"
          "coordinates": [ 4.89994, 52.37815]
        }
      }
    }
  }
}
```



# Data Visualization with Kibana

---

## **What is Nearly-Real-time analysis?**

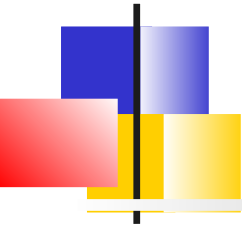
Discover, Visualize and Dashboards

Basic visualisation

Time series

Machine Learning

# Introduction



---

Kibana is an analytics and visualization platform powered by Elasticsearch

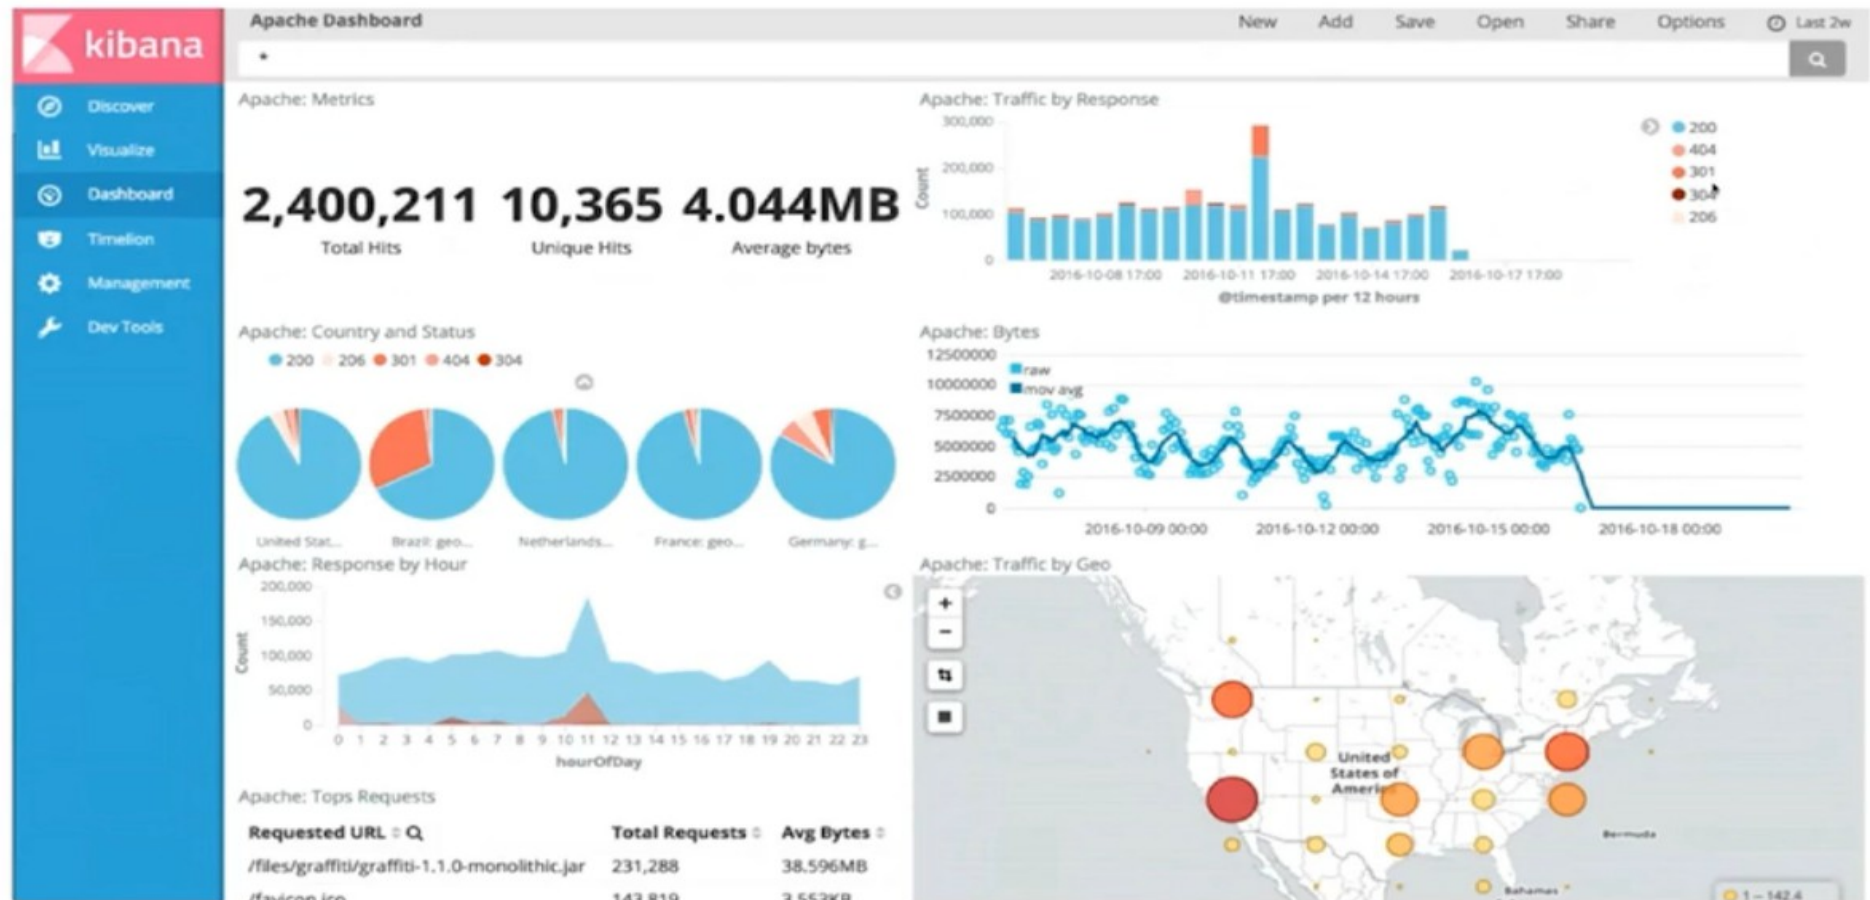
It is able to search data stored in Elasticsearch indexes

It offers a web interface for creating dynamic dashboards displaying the results of queries in real time

It runs on Node.js



# Dashboard Kibana





# Final users and use cases

---

Kibana to destinations of all profiles: Business, Operations, Developer

Once data is ingested into Elasticsearch indexes, it allows:

- Explore the data to better understand it
- Create visualizations and include them in dashboards
- Create presentations to facilitate meetings
- Share via email, web page or PDF documents
- Apply Machine Learning models to detect anomalies or anticipate the future
- Generate graphs visualizing the relationships between your data
- Manage Elasticsearch indexes
- Generate alerts about your data
- Organize Kibana assets into "Space" related to ES security and permissions model



# Data Visualization with Kibana

---

What is Nearly-Real-time analysis?

**Discover, Visualize and Dashboards**

Basic visualisation

Time series

Location analysis

Data relationships with graph

Machine Learning



# Analyze and visualize your data

---

In the context of the Search Engine features of ELK, Kibana can be used :

- Adjust our ELS queries with the ***Dev Console***
- Understand your Data with the menu ***Discover***
- Analyze your data with ***Visualizations*** and ***Dashboards***

# Steps of developing a dashboard



1. Defining an ***Index Pattern/DataView*** and its timestamp field
2. Explore data and optionally save queries
3. Create visualizations, aggregations, times, maps
4. Arrange them on a dashboard
5. Share the Dashboard URL



# Index Pattern creation

---

To exploit data, it is necessary to define ***index pattern***<sup>1</sup>, i.e. collection of indexes respecting a naming rule:

- ex: logstash-\*
- You must also indicate the *timestamp* field of these indexes

Then, Kibana loads the mapping information of these indexes and clearly indicates their type, whether they can be used for search or for aggregation.

# Discover Menu

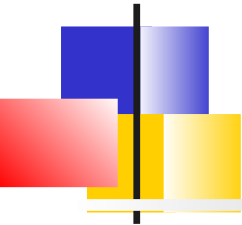
The image shows the Kibana Discover interface with several components labeled:

- Index Pattern:** Points to the `logstash-*` dropdown menu.
- Query bar:** Points to the search bar containing an asterisk `*`.
- Time Picker:** Points to the time range selector showing `May 17th 2015, 04:00:41.685 to May 20th 2015, 18:32:51.964`.
- Toolbar:** Points to the top right area containing buttons for `New`, `Save`, `Open`, `Share`, and a search icon.
- Side Navigation:** Points to the left sidebar menu with options: `Discover`, `Visualize`, `Dashboard`, `Timelion`, `Management`, and `Dev Tools`.
- Histogram:** Points to the bar chart showing the distribution of data over time, with the x-axis labeled `utc_time per hour` and the y-axis labeled `Count`.
- Document Table:** Points to the table of search results below the histogram.

The **Document Table** displays the following data:

Time	_source
May 18th 2015, 02:03:25.877	<code>@timestamp:</code> May 18th 2015, 02:03:25.877 <code>ip:</code> 185.124.182.126 <code>extension:</code> gif <code>response:</code> 404 <code>geo.coordinates:</code> { "lat": 36.518375, "lon": -86.05828083 } <code>geo.src:</code> PH <code>geo.dest:</code> MM <code>geo.srcdest:</code> PH:MM <code>@tags:</code> success, info <code>utc_time:</code> May 18th 2015, 02:03:25.877 <code>referer:</code> http://twitter.com/error/will
May 18th 2015, 05:28:25.013	<code>@timestamp:</code> May 18th 2015, 05:28:25.013 <code>ip:</code> 79.1.14.87 <code>extension:</code> gif <code>response:</code> 200 <code>geo.coordinates:</code> { "lat": 35.16531472, "lon": -107.9006142 } <code>geo.src:</code> GN <code>geo.dest:</code> US <code>geo.srcdest:</code> GN:US <code>@tags:</code> success, info <code>utc_time:</code> May 18th 2015, 05:28:25.013 <code>referer:</code> http://www.slate.com/warning/

# Time window



---

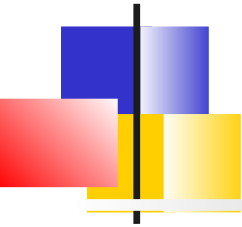
If a timestamp field exists in the index:

- The top of the page displays the distribution of documents over a period (by default 15 min)
- Time window can be changed

This time window is present in the majority of Kibana pages



# Search



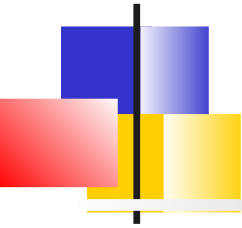
Search criteria can be entered in the query bar. It could be:

- A simple text
- A Lucene query
- A DSL request with JSON.

The result updates the whole page and only the first 500 documents are returned in reverse chronological order

- Field filters can be specified
- Search can be saved with a name
- The index pattern can be modified
- Search can automatically refresh every X times

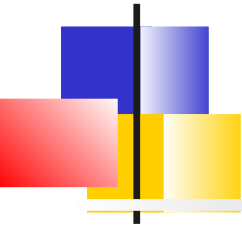
# Filters



Several buttons are available on filters:

- Activation deactivation
- Pin: Filter is retained even on a context switch
- Toggle: Positive or Negative Filter
- Deletion
- Edit: One can then work with DSL syntax and create filter combinations

# Document viewing



---

The list displays by default 500 documents  
(property discover:sampleSize)

It is possible to choose the sorting criterion

Add/Remove Fields

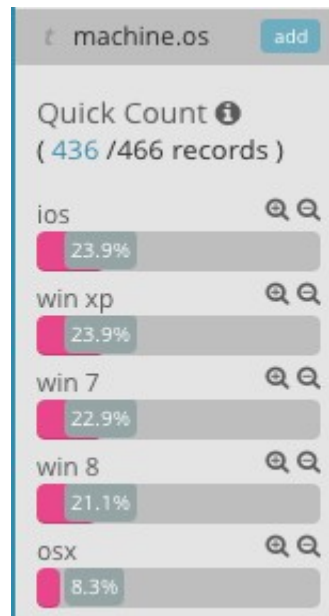
Statistics are also available



# Field value statistics

---

Field value statistics show how many documents have a particular value in a field (facetting)





# Data Visualization with Kibana

---

What is Nearly-Real-time analysis?  
Discover, Visualize and Dashboards

**Basic visualisation**

Time series

Machine Learning

# Visualizations



The visualizations are based on the aggregation capabilities of ELS, we can create graphs that show trends, peaks, ...

Visualizations can then be grouped into dashboards

Different editors for visualization are available:

- Lens: Drag and Drop
- Classical Visualizations based on aggregations
- Maps: Geolocation
- TSVB et Timelion: Suitable for time series
- Custom visualizations: Vega



# Classical visualization

---

**Area** : View total contributions from multiple series

**Table de données** : Tabular display of aggregated data

**Line** : Compare different series

**Metric**: Display a single metric

**Gauge/Goal**: A single value with ranges of good/bad values relative to a target

**Pie** : Pie .

**Heat map** : Heat map

**Nuage de tags** : The most important tags have the largest font

**Horizontal / Vertical bar** : Histogram allowing the comparison of series



# Steps of creation

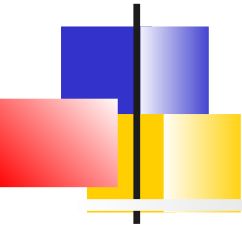
---

With the classical editor steps are

1. Choose a chart type
  2. Specify search criteria or use a saved search
  3. Choose the aggregation calculation for the Y axis (count, average, sum, min, ...)
  4. Choose the data grouping criterion (bucket expression)  
(Histogram of date, Interval, Terms, Filters, Significant terms, ...)
  5. Optionally define sub-aggregations
- Lens allows you to create your visualization more intuitively via Drag & Drop



# Bucket expressions



**Histogramme Date** : Un histogramme de date est construit sur un champ entier pour lequel on a précisé une fenêtre temporel

**Intervalle** : Entier, date ou IPV4

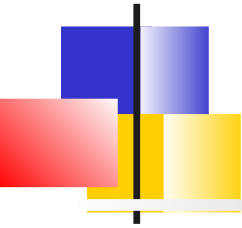
**Terme** : Permet de présente les n meilleurs (ou plus basse) pertinence ordonné par la valeur agrégée

**Filtres** : Il est possible d'ajouter des filtres de données à ce niveau

**Significant Terms** : Agrégation *significant terms*

**Geohash** : Agrégation sur les coordonnées géographiques (Data Table et carte)

# Options



Lors de plusieurs agrégations en Y, il est possible de spécifier leur mode d'affichage

- **Stacked** : Empile les agrégations les unes sur les autres.
- **Overlap** : Superposition avec transparence
- **Wiggle** : Ombrage
- **Percentage** : Chaque agrégation comme une portion du total
- **Silhouette** : Chaque agrégation comme variance d'une ligne centrale
- **Grouped** : Groupe les résultats horizontalement par les sous-agrégations (Grapique barre) .

# Dashboard



---

A dashboard groups a set of saved visualizations into a web page

It is possible to rearrange and resize visualizations

It is possible to share the dashboards by a simple link

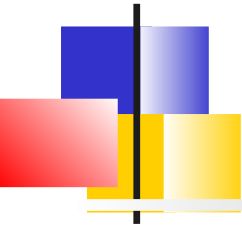


# Data Visualization with Kibana

---

What is Nearly-Real-time analysis?  
Discover, Visualize and Dashboards  
Basic visualisation  
**Time series**  
Machine Learning

# Introduction



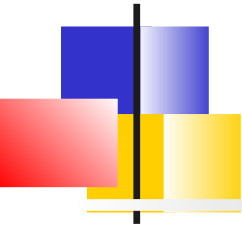
***Timelion*** is a visualizer dedicated to temporal data that allows to combine completely independent data sources in the same graph.

It is driven by a simple expression language to retrieve data and perform calculations.

With timelion, it is possible to answer this type of question:

What percentage of the Japanese population has visited my site?

# Éléments de syntaxe



The timelion interface consists of:

- An expression input box
- A viewing area

Each timelion expression starts with a function identifying a data source.  
For ElasticSearch:

`es(*)`

2 data sources can be drawn side by side using comma (,);.

- `.es(*),.es(metric=cardinality:user)`

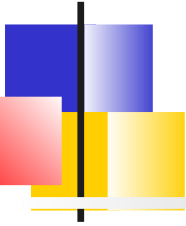
They can be combined via functions:

- `.es(*).divide(es(metric=cardinality:user))`

Other data sources can be used (Worldbank's Data API for example)

`.wbi(MZ) // Population of Mozambique`

# Online help and documentation



Documentation is part of timelion

It is available through the Docs link

In addition, the completion is very efficient

The screenshot shows the Kibana web interface. On the left is a sidebar with navigation links: Discover, Visualize, Dashboard, Timelion (highlighted), Management, and Dev Tools. The main area displays the documentation for the `.es()` function, which is used to pull data from an Elasticsearch instance. At the top of the main area, there is a search bar with `.es|` entered and an 'auto' dropdown. Below the search bar, the title `.es()` Pull data from an elasticsearch instance (Data Source) is shown. A table provides details about the function's arguments.

Argument Name	Accepted Types	Information
q	string, null	Query in lucene query string syntax
metric	string, null	An elasticsearch single value metric agg, eg avg, sum, min, max or cardinality, followed by a field. Eg "sum:bytes", or just "count"
split	string, null	An elasticsearch field to split the series on and a limit. Eg, "hostname:10" to get the top 10 hostnames
index	string, null	Index to query, wildcards accepted
timefield	string, null	Field of type "date" to use for x-axis
kibana	boolean, null	Respect filters on Kibana dashboards. Only has an effect when using on Kibana dashboards



# Examples

---

```
# Real-time display of the average of the  
# user CPU usage percentage from  
# of an elastic search metricbeats index  
.es(index=metricbeat-*, timefield='@timestamp',  
metric='avg:system.cpu.user.pct')
```

```
# The same thing by adding a series displaying  
# the data for an offset of 1h  
.es(index=metricbeat-*, timefield='@timestamp',  
metric='avg:system.cpu.user.pct'), .es(offset=-  
1h,index=metricbeat-*, timefield='@timestamp',  
metric='avg:system.cpu.user.pct')
```





# Data Visualization with Kibana

---

What is Nearly-Real-time analysis?  
Discover, Visualize and Dashboards

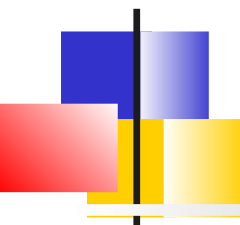
Basic visualisation

Time series

Location analysis

**Machine Learning**

# Introduction



Elastic's ML offering, although applicable to many areas, focuses on IT monitoring:

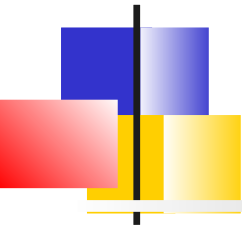
- Monitoring: Detection of failure, malfunction
- Security: Intrusion, exfiltration detection

In IT monitoring:

We generate a lot of data but we don't have time to look at it and analyze it

=> Automation of these analyzes by systems that can learn on their own

# Anomalies



---

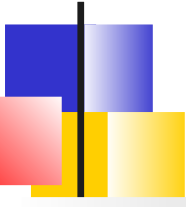
The majority of monitoring requirements are variations on the theme:

*Find something that is different from the ordinary*

The analysis processes then attempt to detect anomalies in relation to the observation of the history

=> This requires the history to be large enough so that the ML can refine its model

# Objectives of a ML solution



**Speed:** Notification of a failure, a breach or any other significant anomaly proactively and as quickly as possible in order to mitigate it.

**Scalability:** Algorithms must be able to scale linearly with data.

**Performance:** Using modest hardware rather than supercomputers

**Applicability:** Consideration of the diversity of data in computing environments.

**Adaptability:** Constantly changing computing environments can quickly make a static algorithm fragile.

**Accuracy:** Do not generate false alarms

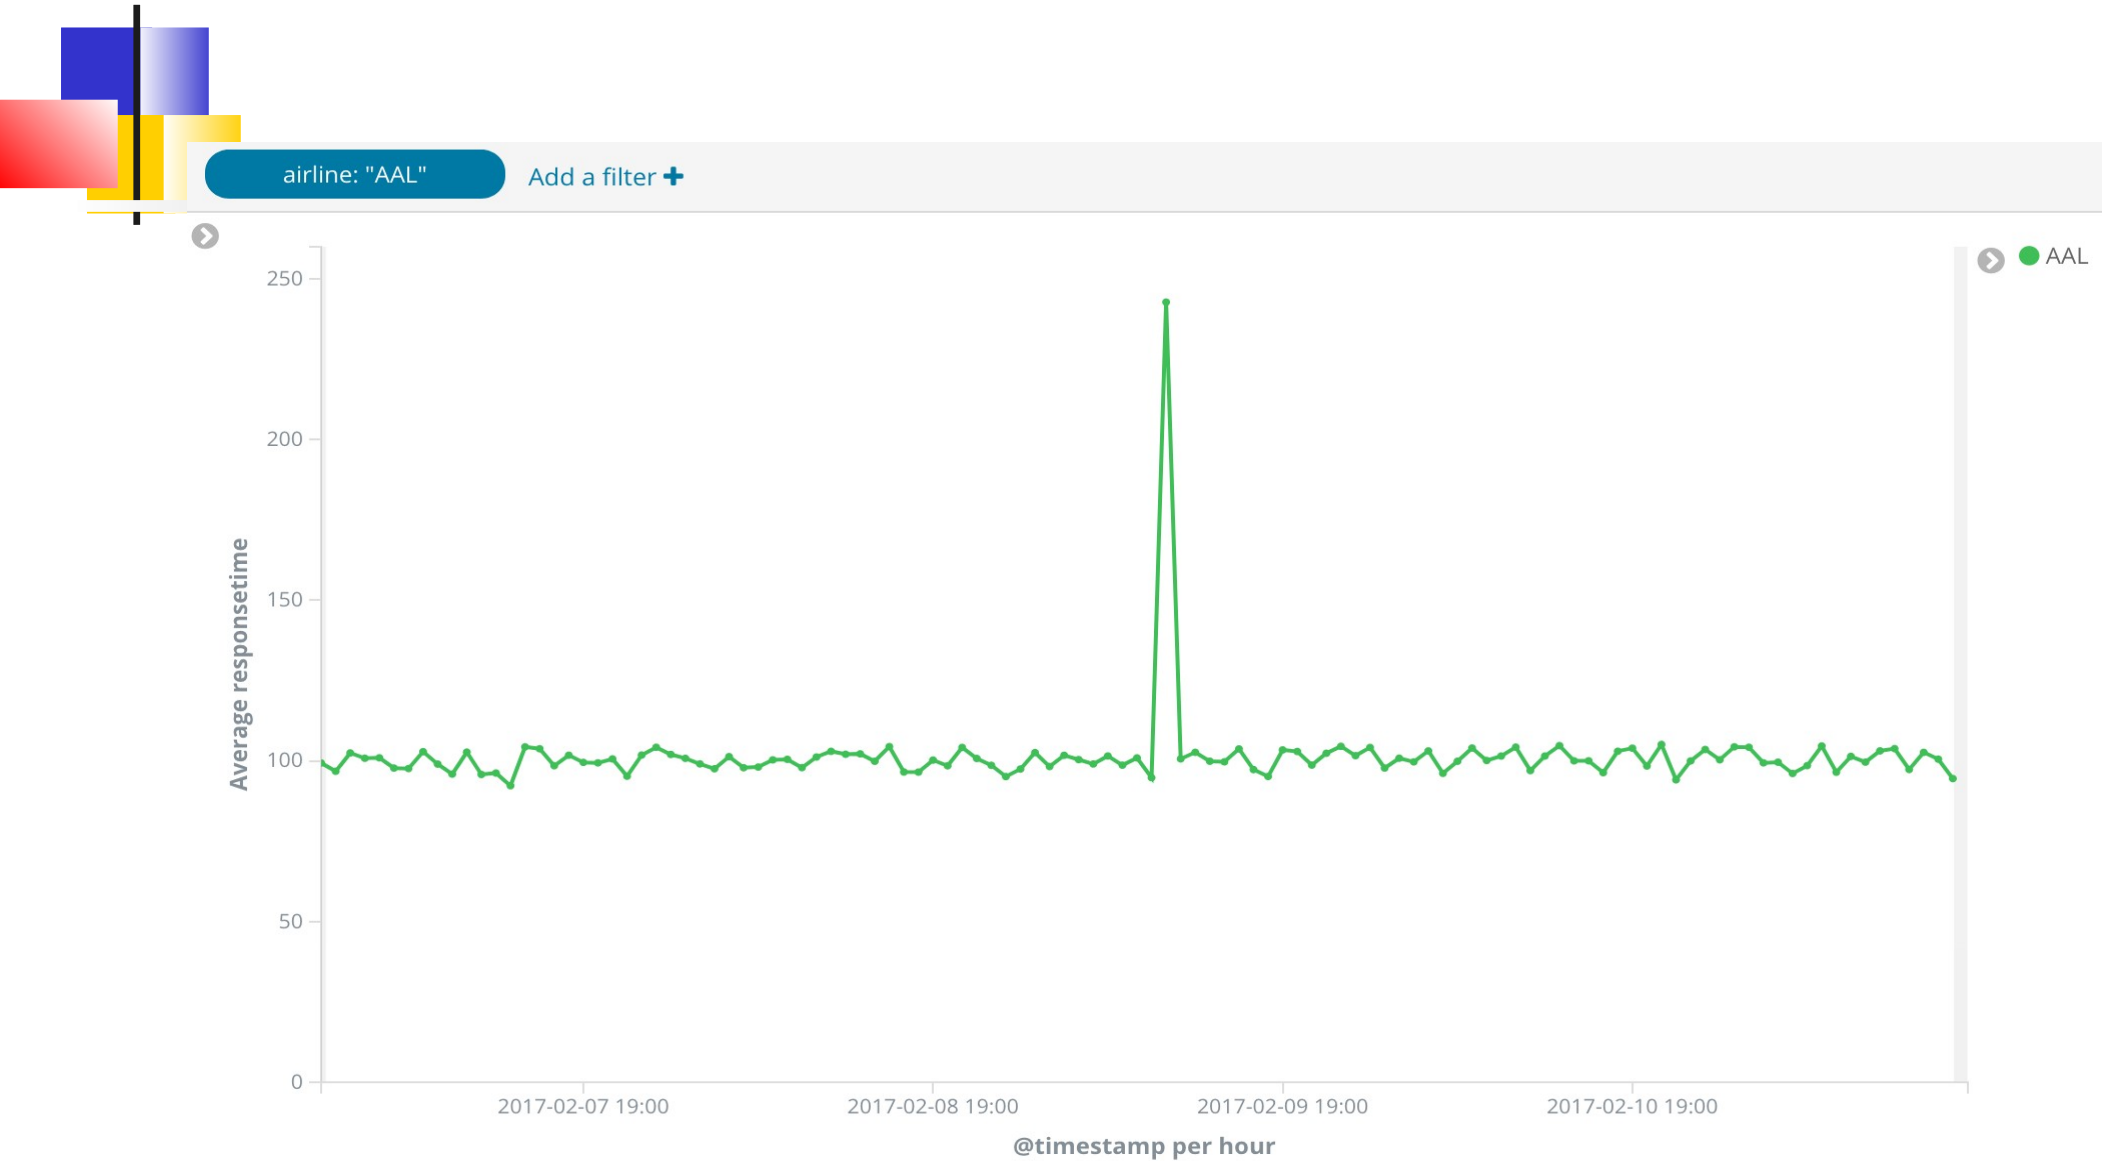
# What is an anomaly ?



## 2 definitions of an anomaly:

- Something is unusual if its behavior has deviated significantly from an established pattern based on its past history
- Something is unusual if some of its characteristics are significantly different from the same characteristics of other members of a set or population

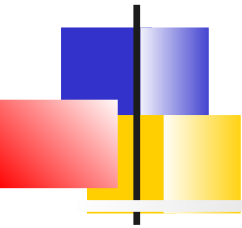
# Ex : Past history



# Ex : Population



# Unsupervised learning



In ML, unsupervised learning requires no human intervention for setting up and refining the model.

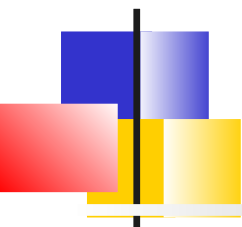
- => The model is built, refined, modified according to the data presented to it
- => The more data there is, the more accurate the model becomes

Elastic Stack offers 2 types of unsupervised analysis

- Anomaly detection: Based on temporal data
- Abnormal value detection: Based on value densities



# Supervised learning



Supervised learning requires training data.

ELK-ML offers supervised analysis types:

- **Classification**: Event Classification  
Ex: Normal or malicious requests
- **Regression**: Predicting Numerical Values  
Ex: response for a web request.

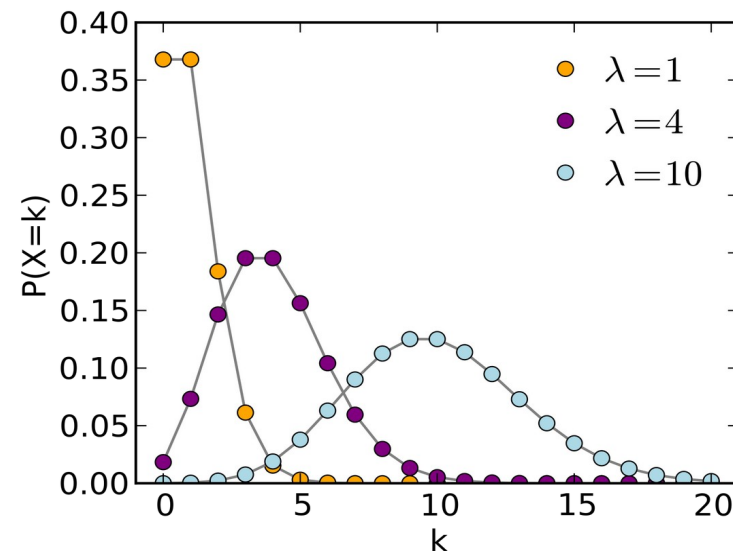
# Ex : Poisson Distribution

Statistical models are used to model the behavior of a system.

For example, the Poisson distribution makes it possible to model events such as

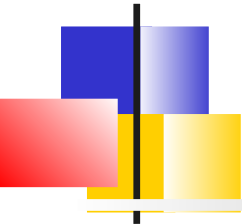
- The number of meteorites over a meter in diameter that hit Earth each year
- The number of patients arriving at an emergency room between 10:00 a.m. and 11:00 p.m.
- The number of photons hitting a detector in a particular time interval

$$p(k) = \mathbb{P}(X = k) = \frac{\lambda^k}{k!} e^{-\lambda}$$



If a certain type of event occurs on average 4 times per minute,  
to study the number of events occurring in a 10-minute time frame,  
we choose as a model a Poisson distribution with parameter  $\lambda = 10 \times 4 = 40$

# Modelisation process



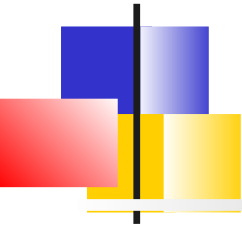
The machine learning process selects from its observations the appropriate probability model (Poisson, Gaussian, log-normal, etc.) and its coefficients

Bayesian techniques are used to assess the probabilities of model values given the observed data set and allow the results to be tempered based on the amount of information visualized.

The modeling performed is continuous, so new information is considered along with old, with an exponential weighting of fresher information

Modeling can be stopped and then restarted later, hence the need to persist the model.

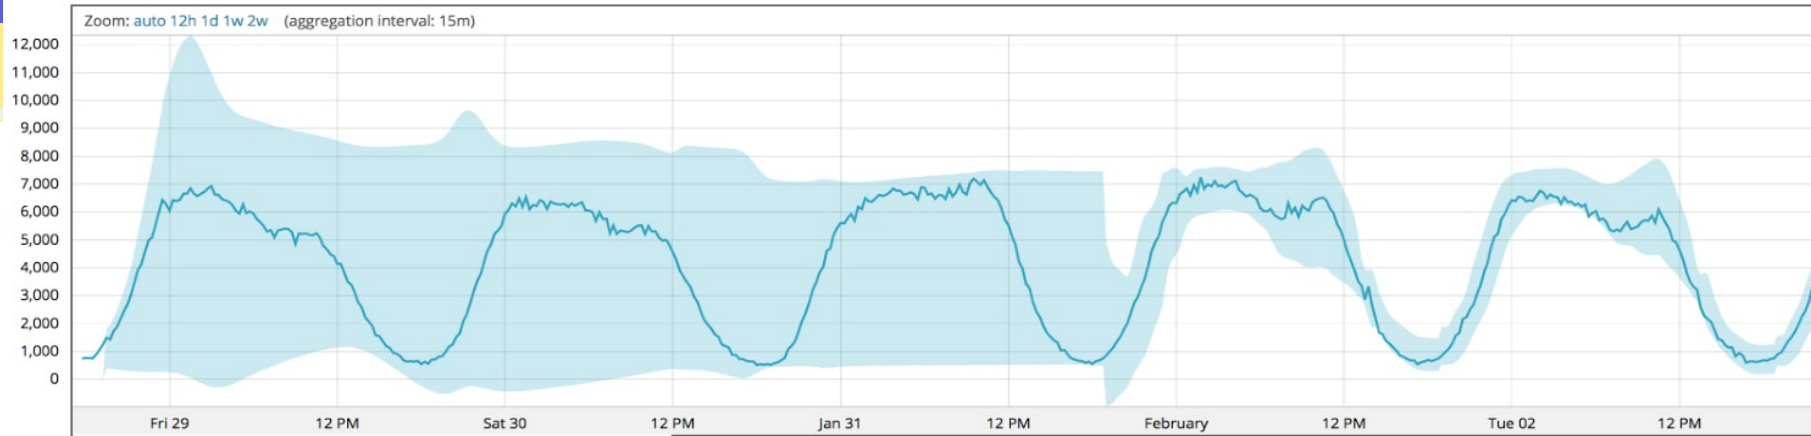
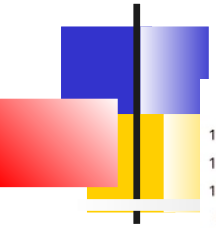
# Cycles



Another important aspect of modeling real data is to take into account the important harmonic trends that occur naturally.

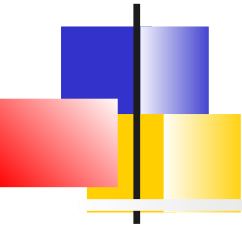
ML automatically finds salient trends in the data (linear growth, cyclic harmonics, etc.) and factors them

# Example



The periodic daily cycle is learned and then factored.  
The model prediction adjusts after the automatic detection of three successive iterations of this cycle.

# Anomaly Score



---

Probability models are used to calculate the probability rate that a measurement takes on a certain value.

Probability rates oscillating between 0 and 1 are normalized to **anomaly scores** that oscillate between 0 and 100

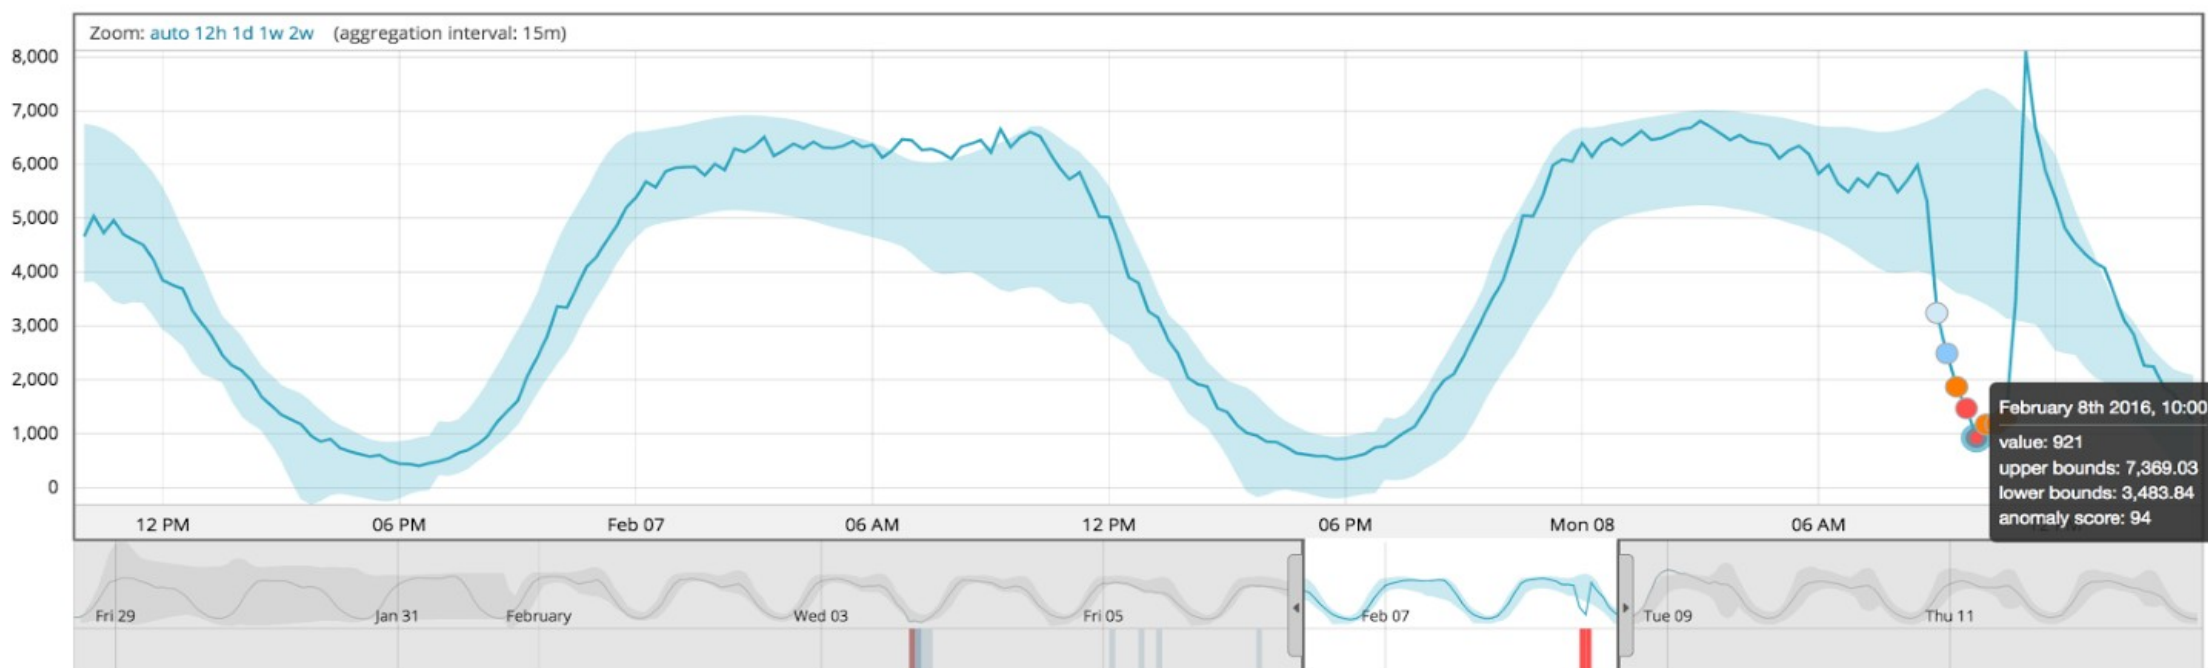
Severity levels, alerts are associated with anomaly score thresholds

# Example

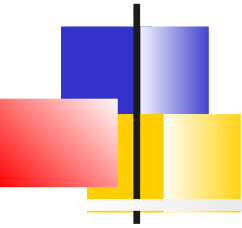
The probability of the value 921 was  $6.3634e-7$

This gives an anomaly score of 94

Which gives a critical anomaly ( $> 75$ )



# Forecasts



---

Once ML has created its behavioral model, it uses it to extrapolate future behavior:

- Estimate a value at a future date
- Probability that a value reaches a threshold

Each forecast has a unique id, duration, expiry time

Forecasts cannot be made on all types of values

They are limited to 3 for the same job

They require a lot of memory



# Example

